

Chapter_3_A_Tour_of_Machine_Learning_Classifiers_Using_Scikit-Learn

March 17, 2024

0.1 Implementing Perceptron using scikit-learn on iris dataset

```
[34]: from sklearn import datasets
import numpy as np
iris = datasets.load_iris()
X = iris.data[:, [2, 3]]
y = iris.target
print('Class labels:', np.unique(y))
```

Class labels: [0 1 2]

```
[35]: # iris
```

```
[36]: from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3,
↪random_state=1, stratify=y)
```

```
[37]: print('Labels counts in y:', np.bincount(y))
print('Labels counts in y_train:', np.bincount(y_train))
print('Labels counts in y_test:', np.bincount(y_test))
```

Labels counts in y: [50 50 50]

Labels counts in y_train: [35 35 35]

Labels counts in y_test: [15 15 15]

```
[38]: from sklearn.preprocessing import StandardScaler
sc = StandardScaler()
sc.fit(X_train)
X_train_std = sc.transform(X_train)
X_test_std = sc.transform(X_test)
```

```
[39]: from sklearn.linear_model import Perceptron
ppn = Perceptron(max_iter=40, eta0=0.1, random_state=1)
ppn.fit(X_train_std, y_train)
```

```
[39]: Perceptron(eta0=0.1, max_iter=40, random_state=1)
```

```
[40]: y_pred = ppn.predict(X_test_std)
print('Misclassified samples: %d' % (y_test != y_pred).sum())
```

Misclassified samples: 1

```
[41]: from sklearn.metrics import accuracy_score
print('Accuracy: %.2f' % accuracy_score(y_test, y_pred))
```

Accuracy: 0.98

```
[42]: print('Accuracy: %.2f' % ppn.score(X_test_std, y_test))
```


Accuracy: 0.98

```
[43]: from matplotlib.colors import ListedColormap
import matplotlib.pyplot as plt
def plot_decision_regions(X, y, classifier, test_idx=None, resolution=0.02):
    # setup marker generator and color map
    markers = ('s', 'x', 'o', '^', 'v')
    colors = ('red', 'blue', 'lightgreen', 'gray', 'cyan')
    cmap = ListedColormap(colors[:len(np.unique(y))])
    # plot the decision surface
    x1_min, x1_max = X[:, 0].min() - 1, X[:, 0].max() + 1
    x2_min, x2_max = X[:, 1].min() - 1, X[:, 1].max() + 1
    xx1, xx2 = np.meshgrid(np.arange(x1_min, x1_max, resolution),
                           np.arange(x2_min, x2_max, resolution))
    Z = classifier.predict(np.array([xx1.ravel(), xx2.ravel()]).T)
    Z = Z.reshape(xx1.shape)
    plt.contourf(xx1, xx2, Z, alpha=0.3, cmap=cmap)
    plt.xlim(xx1.min(), xx1.max())
    plt.ylim(xx2.min(), xx2.max())
    for idx, cl in enumerate(np.unique(y)):
        plt.scatter(x=X[y == cl, 0], y=X[y == cl, 1], alpha=0.8,
                    c=colors[idx], marker=markers[idx], label=cl, edgecolor='black')
    # highlight test samples
    if test_idx:
        # plot all samples
        X_test, y_test = X[test_idx, :], y[test_idx]
        plt.scatter(X_test[:, 0], X_test[:, 1], edgecolor='black', alpha=1.
                    ↪0, linewidth=1, marker='o', s=100, label='test set')
```

```
[44]: X_combined_std = np.vstack((X_train_std, X_test_std))
y_combined = np.hstack((y_train, y_test))
plot_decision_regions(X=X_combined_std, y=y_combined, classifier=ppn, test_idx=range(105, ↪
↪150))
plt.xlabel('petal length [standardized]')
plt.ylabel('petal width [standardized]')
plt.legend(loc='upper left')
plt.show()
```

C:\Users\ankit19.gupta\AppData\Local\Temp\ipykernel_15304\1174551119.py:19:
 UserWarning: You passed a edgecolor/edgecolors ('black') for an unfilled marker ('x'). Matplotlib is ignoring the edgecolor in favor of the facecolor. This behavior may change in the future.

```
plt.scatter(x=X[y == c1, 0], y=X[y == c1, 1],alpha=0.8,
c=colors[idx],marker=markers[idx], label=c1,edgecolor='black')
```



[45]: *## decision boundary is not perfectly separable because data classes are not linearly separable. So, use logistic algorithm*

[46]: `help(Perceptron)`

Help on class Perceptron in module sklearn.linear_model._perceptron:

```
class Perceptron(sklearn.linear_model._stochastic_gradient.BaseSGDClassifier)
|   Perceptron(*, penalty=None, alpha=0.0001, l1_ratio=0.15, fit_intercept=True,
|   max_iter=1000, tol=0.001, shuffle=True, verbose=0, eta0=1.0, n_jobs=None,
|   random_state=0, early_stopping=False, validation_fraction=0.1,
|   n_iter_no_change=5, class_weight=None, warm_start=False)
|
|   Linear perceptron classifier.
```

```

| The implementation is a wrapper around
:~sklearn.linear_model.SGDClassifier`
| by fixing the `loss` and `learning_rate` parameters as::
|
|     SGDClassifier(loss="perceptron", learning_rate="constant")
|
| Other available parameters are described below and are forwarded to
| :~sklearn.linear_model.SGDClassifier`.
|
| Read more in the :ref:`User Guide <perceptron>`.
|
| Parameters
| -----
|
| penalty : {'l2', 'l1', 'elasticnet'}, default=None
|     The penalty (aka regularization term) to be used.
|
| alpha : float, default=0.0001
|     Constant that multiplies the regularization term if regularization is
|     used.
|
| l1_ratio : float, default=0.15
|     The Elastic Net mixing parameter, with `0 <= l1_ratio <= 1`.
|     `l1_ratio=0` corresponds to L2 penalty, `l1_ratio=1` to L1.
|     Only used if `penalty='elasticnet'`.
|
|     .. versionadded:: 0.24
|
| fit_intercept : bool, default=True
|     Whether the intercept should be estimated or not. If False, the
|     data is assumed to be already centered.
|
| max_iter : int, default=1000
|     The maximum number of passes over the training data (aka epochs).
|     It only impacts the behavior in the ``fit`` method, and not the
|     :meth:`partial_fit` method.
|
|     .. versionadded:: 0.19
|
| tol : float or None, default=1e-3
|     The stopping criterion. If it is not None, the iterations will stop
|     when (loss > previous_loss - tol).
|
|     .. versionadded:: 0.19
|
| shuffle : bool, default=True
|     Whether or not the training data should be shuffled after each epoch.

```

```

| verbose : int, default=0
|     The verbosity level.
|
| eta0 : float, default=1
|     Constant by which the updates are multiplied.
|
| n_jobs : int, default=None
|     The number of CPUs to use to do the OVA (One Versus All, for
|     multi-class problems) computation.
|     ``None`` means 1 unless in a :obj:`joblib.parallel_backend` context.
|     ``-1`` means using all processors. See :term:`Glossary <n_jobs>`
|     for more details.
|
| random_state : int, RandomState instance or None, default=0
|     Used to shuffle the training data, when ``shuffle`` is set to
|     ``True``. Pass an int for reproducible output across multiple
|     function calls.
|     See :term:`Glossary <random_state>`.
|
| early_stopping : bool, default=False
|     Whether to use early stopping to terminate training when validation
|     score is not improving. If set to True, it will automatically set aside
|     a stratified fraction of training data as validation and terminate
|     training when validation score is not improving by at least `tol` for
|     `n_iter_no_change` consecutive epochs.
|
|     .. versionadded:: 0.20
|
| validation_fraction : float, default=0.1
|     The proportion of training data to set aside as validation set for
|     early stopping. Must be between 0 and 1.
|     Only used if early_stopping is True.
|
|     .. versionadded:: 0.20
|
| n_iter_no_change : int, default=5
|     Number of iterations with no improvement to wait before early stopping.
|
|     .. versionadded:: 0.20
|
| class_weight : dict, {class_label: weight} or "balanced", default=None
|     Preset for the class_weight fit parameter.
|
|     Weights associated with classes. If not given, all classes
|     are supposed to have weight one.
|
|     The "balanced" mode uses the values of y to automatically adjust
|     weights inversely proportional to class frequencies in the input data

```

```

|         as ``n_samples / (n_classes * np.bincount(y))``.
|
| warm_start : bool, default=False
|     When set to True, reuse the solution of the previous call to fit as
|     initialization, otherwise, just erase the previous solution. See
|     :term:`the Glossary <warm_start>`.
|
| Attributes
| -----
| classes_ : ndarray of shape (n_classes,)
|     The unique classes labels.
|
| coef_ : ndarray of shape (1, n_features) if n_classes == 2 else
(n_classes, n_features)
|     Weights assigned to the features.
|
| intercept_ : ndarray of shape (1,) if n_classes == 2 else (n_classes,)
|     Constants in decision function.
|
| loss_function_ : concrete LossFunction
|     The function that determines the loss, or difference between the
|     output of the algorithm and the target values.
|
| n_features_in_ : int
|     Number of features seen during :term:`fit`.
|
|     .. versionadded:: 0.24
|
| feature_names_in_ : ndarray of shape (`n_features_in_`,)
|     Names of features seen during :term:`fit`. Defined only when `X`
|     has feature names that are all strings.
|
|     .. versionadded:: 1.0
|
| n_iter_ : int
|     The actual number of iterations to reach the stopping criterion.
|     For multiclass fits, it is the maximum over every binary fit.
|
| t_ : int
|     Number of weight updates performed during training.
|     Same as `` (n_iter_ * n_samples + 1) ``.
|
| See Also
| -----
| sklearn.linear_model.SGDClassifier : Linear classifiers
|     (SVM, logistic regression, etc.) with SGD training.
|
| Notes

```

```

| -----
| ``Perceptron`` is a classification algorithm which shares the same
| underlying implementation with ``SGDClassifier``. In fact,
| ``Perceptron()`` is equivalent to ``SGDClassifier(loss="perceptron",
| eta0=1, learning_rate="constant", penalty=None)`.
|
| References
| -----
| https://en.wikipedia.org/wiki/Perceptron and references therein.
|
| Examples
| -----
| >>> from sklearn.datasets import load_digits
| >>> from sklearn.linear_model import Perceptron
| >>> X, y = load_digits(return_X_y=True)
| >>> clf = Perceptron(tol=1e-3, random_state=0)
| >>> clf.fit(X, y)
| Perceptron()
| >>> clf.score(X, y)
| 0.939...
|
| Method resolution order:
|     Perceptron
|     sklearn.linear_model._stochastic_gradient.BaseSGDClassifier
|     sklearn.linear_model._base.LinearClassifierMixin
|     sklearn.base.ClassifierMixin
|     sklearn.linear_model._stochastic_gradient.BaseSGD
|     sklearn.linear_model._base.SparseCoefMixin
|     sklearn.base.BaseEstimator
|     sklearn.utils._estimator_html_repr._HTMLDocumentationLinkMixin
|     sklearn.utils._metadata_requests._MetadataRequester
|     builtins.object
|
| Methods defined here:
|
|     __init__(self, *, penalty=None, alpha=0.0001, l1_ratio=0.15,
| fit_intercept=True, max_iter=1000, tol=0.001, shuffle=True, verbose=0, eta0=1.0,
| n_jobs=None, random_state=0, early_stopping=False, validation_fraction=0.1,
| n_iter_no_change=5, class_weight=None, warm_start=False)
|         Initialize self. See help(type(self)) for accurate signature.
|
|     set_fit_request(self: sklearn.linear_model._perceptron.Perceptron, *,
| coef_init: Union[bool, NoneType, str] = '$UNCHANGED$', intercept_init:
| Union[bool, NoneType, str] = '$UNCHANGED$', sample_weight: Union[bool, NoneType,
| str] = '$UNCHANGED$') -> sklearn.linear_model._perceptron.Perceptron
|         Request metadata passed to the ``fit`` method.
|
|     Note that this method is only relevant if

```

```

|     ``enable_metadata_routing=True`` (see :func:`sklearn.set_config`).
|     Please see :ref:`User Guide <metadata_routing>` on how the routing
|     mechanism works.
|
|     The options for each parameter are:
|
|     - ``True``: metadata is requested, and passed to ``fit`` if provided.
The request is ignored if metadata is not provided.
|
|     - ``False``: metadata is not requested and the meta-estimator will not
pass it to ``fit``.
|
|     - ``None``: metadata is not requested, and the meta-estimator will raise
an error if the user provides it.
|
|     - ``str``: metadata should be passed to the meta-estimator with this
given alias instead of the original name.
|
|     The default (``sklearn.utils.metadata_routing.UNCHANGED``) retains the
|     existing request. This allows you to change the request for some
|     parameters and not others.
|
|     .. versionadded:: 1.3
|
|     .. note::
|         This method is only relevant if this estimator is used as a
|         sub-estimator of a meta-estimator, e.g. used inside a
|         :class:`~sklearn.pipeline.Pipeline`. Otherwise it has no effect.
|
|     Parameters
|     -----
|
|     coef_init : str, True, False, or None,
default=sklearn.utils.metadata_routing.UNCHANGED
|         Metadata routing for ``coef_init`` parameter in ``fit``.
|
|     intercept_init : str, True, False, or None,
default=sklearn.utils.metadata_routing.UNCHANGED
|         Metadata routing for ``intercept_init`` parameter in ``fit``.
|
|     sample_weight : str, True, False, or None,
default=sklearn.utils.metadata_routing.UNCHANGED
|         Metadata routing for ``sample_weight`` parameter in ``fit``.
|
|     Returns
|     -----
|
|     self : object
|         The updated object.
|

```



```

| set_partial_fit_request(self: sklearn.linear_model._perceptron.Perceptron,
*, classes: Union[bool, NoneType, str] = '$UNCHANGED$', sample_weight:
Union[bool, NoneType, str] = '$UNCHANGED$') ->
sklearn.linear_model._perceptron.Perceptron
|     Request metadata passed to the ``partial_fit`` method.
|
|     Note that this method is only relevant if
|     ``enable_metadata_routing=True`` (see :func:`sklearn.set_config`).
|     Please see :ref:`User Guide <metadata_routing>` on how the routing
|     mechanism works.
|
|     The options for each parameter are:
|
|     - ``True``: metadata is requested, and passed to ``partial_fit`` if
provided. The request is ignored if metadata is not provided.
|
|     - ``False``: metadata is not requested and the meta-estimator will not
pass it to ``partial_fit``.
|
|     - ``None``: metadata is not requested, and the meta-estimator will raise
an error if the user provides it.
|
|     - ``str``: metadata should be passed to the meta-estimator with this
given alias instead of the original name.
|
|     The default (``sklearn.utils.metadata_routing.UNCHANGED``) retains the
|     existing request. This allows you to change the request for some
|     parameters and not others.
|
|     .. versionadded:: 1.3
|
|     .. note::
|         This method is only relevant if this estimator is used as a
|         sub-estimator of a meta-estimator, e.g. used inside a
|         :class:`~sklearn.pipeline.Pipeline`. Otherwise it has no effect.
|
|     Parameters
|     -----
|
|     classes : str, True, False, or None,
default=sklearn.utils.metadata_routing.UNCHANGED
|         Metadata routing for ``classes`` parameter in ``partial_fit``.
|
|     sample_weight : str, True, False, or None,
default=sklearn.utils.metadata_routing.UNCHANGED
|         Metadata routing for ``sample_weight`` parameter in ``partial_fit``.
|
|     Returns
|     -----

```

```

|     self : object
|         The updated object.
|
|     set_score_request(self: sklearn.linear_model._perceptron.Perceptron, *,
sample_weight: Union[bool, NoneType, str] = '$UNCHANGED$') ->
sklearn.linear_model._perceptron.Perceptron
|         Request metadata passed to the ``score`` method.
|
|         Note that this method is only relevant if
|         ``enable_metadata_routing=True`` (see :func:`sklearn.set_config`).
|         Please see :ref:`User Guide <metadata_routing>` on how the routing
|         mechanism works.
|
|         The options for each parameter are:
|
|         - ``True``: metadata is requested, and passed to ``score`` if provided.
The request is ignored if metadata is not provided.
|
|         - ``False``: metadata is not requested and the meta-estimator will not
pass it to ``score``.
|
|         - ``None``: metadata is not requested, and the meta-estimator will raise
an error if the user provides it.
|
|         - ``str``: metadata should be passed to the meta-estimator with this
given alias instead of the original name.
|
|         The default (``sklearn.utils.metadata_routing.UNCHANGED``) retains the
|         existing request. This allows you to change the request for some
|         parameters and not others.
|
|         .. versionadded:: 1.3
|
|         .. note::
|             This method is only relevant if this estimator is used as a
|             sub-estimator of a meta-estimator, e.g. used inside a
|             :class:`~sklearn.pipeline.Pipeline`. Otherwise it has no effect.
|
|         Parameters
|         -----
|
|         sample_weight : str, True, False, or None,
default=sklearn.utils.metadata_routing.UNCHANGED
|             Metadata routing for ``sample_weight`` parameter in ``score``.
|
|         Returns
|         -----
|
|         self : object
|             The updated object.

```

```

| -----
| Data and other attributes defined here:
|
| __abstractmethods__ = frozenset()
|
| __annotations__ = {'_parameter_constraints': <class 'dict'>}
| -----
| Methods inherited from
sklearn.linear_model._stochastic_gradient.BaseSGDClassifier:
|
| fit(self, X, y, coef_init=None, intercept_init=None, sample_weight=None)
|     Fit linear model with Stochastic Gradient Descent.
|
|     Parameters
|     -----
|     X : {array-like, sparse matrix}, shape (n_samples, n_features)
|         Training data.
|
|     y : ndarray of shape (n_samples,)
|         Target values.
|
|     coef_init : ndarray of shape (n_classes, n_features), default=None
|         The initial coefficients to warm-start the optimization.
|
|     intercept_init : ndarray of shape (n_classes,), default=None
|         The initial intercept to warm-start the optimization.
|
|     sample_weight : array-like, shape (n_samples,), default=None
|         Weights applied to individual samples.
|         If not provided, uniform weights are assumed. These weights will
|         be multiplied with class_weight (passed through the
|         constructor) if class_weight is specified.
|
|     Returns
|     -----
|     self : object
|         Returns an instance of self.
|
| partial_fit(self, X, y, classes=None, sample_weight=None)
|     Perform one epoch of stochastic gradient descent on given samples.
|
|     Internally, this method uses ``max_iter = 1``. Therefore, it is not
|     guaranteed that a minimum of the cost function is reached after calling
|     it once. Matters such as objective convergence, early stopping, and
|     learning rate adjustments should be handled by the user.

```

```

| Parameters
| -----
| X : {array-like, sparse matrix}, shape (n_samples, n_features)
|     Subset of the training data.
|
| y : ndarray of shape (n_samples,)
|     Subset of the target values.
|
| classes : ndarray of shape (n_classes,), default=None
|     Classes across all calls to partial_fit.
|     Can be obtained by via `np.unique(y_all)`, where y_all is the
|     target vector of the entire dataset.
|     This argument is required for the first call to partial_fit
|     and can be omitted in the subsequent calls.
|     Note that y doesn't need to contain all labels in `classes`.
|
| sample_weight : array-like, shape (n_samples,), default=None
|     Weights applied to individual samples.
|     If not provided, uniform weights are assumed.
|
| Returns
| -----
| self : object
|     Returns an instance of self.
|
| -----
| Data and other attributes inherited from
sklearn.linear_model._stochastic_gradient.BaseSGDClassifier:
|
| loss_functions = {'epsilon_insensitive': (<class 'sklearn.linear_model...
|
| -----
| Methods inherited from sklearn.linear_model._base.LinearClassifierMixin:
|
| decision_function(self, X)
|     Predict confidence scores for samples.
|
|     The confidence score for a sample is proportional to the signed
|     distance of that sample to the hyperplane.
|
| Parameters
| -----
| X : {array-like, sparse matrix} of shape (n_samples, n_features)
|     The data matrix for which we want to get the confidence scores.
|
| Returns
| -----
| scores : ndarray of shape (n_samples,) or (n_samples, n_classes)

```

```

|         Confidence scores per `(n_samples, n_classes)` combination. In the
|         binary case, confidence score for `self.classes_[1]` where >0 means
|         this class would be predicted.
|
| predict(self, X)
|     Predict class labels for samples in X.
|
|     Parameters
|     -----
|     X : {array-like, sparse matrix} of shape (n_samples, n_features)
|         The data matrix for which we want to get the predictions.
|
|     Returns
|     -----
|     y_pred : ndarray of shape (n_samples,)
|         Vector containing the class labels for each sample.
|
| -----
| Methods inherited from sklearn.base.ClassifierMixin:
|
| score(self, X, y, sample_weight=None)
|     Return the mean accuracy on the given test data and labels.
|
|     In multi-label classification, this is the subset accuracy
|     which is a harsh metric since you require for each sample that
|     each label set be correctly predicted.
|
|     Parameters
|     -----
|     X : array-like of shape (n_samples, n_features)
|         Test samples.
|
|     y : array-like of shape (n_samples,) or (n_samples, n_outputs)
|         True labels for `X`.
|
|     sample_weight : array-like of shape (n_samples,), default=None
|         Sample weights.
|
|     Returns
|     -----
|     score : float
|         Mean accuracy of ``self.predict(X)`` w.r.t. `y`.
|
| -----
| Data descriptors inherited from sklearn.base.ClassifierMixin:
|
| __dict__
|     dictionary for instance variables (if defined)

```

```

|  __weakref__
|      list of weak references to the object (if defined)
|
|  -----
|  Readonly properties inherited from
sklearn.linear_model._stochastic_gradient.BaseSGD:
|
|  loss_function_
|
|  -----
|  Methods inherited from sklearn.linear_model._base.SparseCoefMixin:
|
|  densify(self)
|      Convert coefficient matrix to dense array format.
|
|      Converts the ``coef_`` member (back) to a numpy.ndarray. This is the
|      default format of ``coef_`` and is required for fitting, so calling
|      this method is only required on models that have previously been
|      sparsified; otherwise, it is a no-op.
|
|      Returns
|      -----
|      self
|          Fitted estimator.
|
|  sparsify(self)
|      Convert coefficient matrix to sparse format.
|
|      Converts the ``coef_`` member to a scipy.sparse matrix, which for
|      L1-regularized models can be much more memory- and storage-efficient
|      than the usual numpy.ndarray representation.
|
|      The ``intercept_`` member is not converted.
|
|      Returns
|      -----
|      self
|          Fitted estimator.
|
|      Notes
|      -----
|      For non-sparse models, i.e. when there are not many zeros in ``coef_``,
|      this may actually increase memory usage, so use this method with
|      care. A rule of thumb is that the number of zero elements, which can
|      be computed with ``(coef_ == 0).sum()`` , must be more than 50% for this
|      to provide significant benefits.

```

After calling this method, further fitting with the `partial_fit` method (if any) will not work until you call `densify`.

Methods inherited from `sklearn.base.BaseEstimator`:

`__getstate__(self)`

`__repr__(self, N_CHAR_MAX=700)`
Return `repr(self)`.

`__setstate__(self, state)`

`__sklearn_clone__(self)`

`get_params(self, deep=True)`
Get parameters for this estimator.

Parameters

`deep` : bool, default=True
If True, will return the parameters for this estimator and contained subobjects that are estimators.

Returns

`params` : dict
Parameter names mapped to their values.

`set_params(self, **params)`
Set the parameters of this estimator.

The method works on simple estimators as well as on nested objects (such as `:class:`~sklearn.pipeline.Pipeline``). The latter have parameters of the form ``<component>__<parameter>`` so that it's possible to update each component of a nested object.

Parameters

`**params` : dict
Estimator parameters.

Returns

`self` : estimator instance
Estimator instance.

```

| Methods inherited from sklearn.utils._metadata_requests._MetadataRequester:
|
| get_metadata_routing(self)
|     Get metadata routing of this object.
|
|     Please check :ref:`User Guide <metadata_routing>` on how the routing
|     mechanism works.
|
|     Returns
|     -----
|     routing : MetadataRequest
|         A :class:`~sklearn.utils.metadata_routing.MetadataRequest`
encapsulating
|         routing information.
|
| -----
| Class methods inherited from
sklearn.utils._metadata_requests._MetadataRequester:
|
| __init_subclass__(**kwargs) from abc.ABCMeta
|     Set the ``set_{method}_request`` methods.
|
|     This uses PEP-487 [1]_ to set the ``set_{method}_request`` methods. It
|     looks for the information available in the set default values which are
|     set using ``__metadata_request__*`` class attributes, or inferred
|     from method signatures.
|
|     The ``__metadata_request__*`` class attributes are used when a method
|     does not explicitly accept a metadata through its arguments or if the
|     developer would like to specify a request value for those metadata
|     which are different from the default ``None``.
|
|     References
|     -----
|     .. [1] https://www.python.org/dev/peps/pep-0487

```

0.2 Implementing Logistic Regression using scikit-learn on iris dataset

```

[47]: #Now let us simply plot the sigmoid function for some values in the range -7 to 7
      ↪ 7 to
      #see how it looks:
      import matplotlib.pyplot as plt
      import numpy as np
      def sigmoid(z):
          return 1.0 / (1.0 + np.exp(-z))
      z = np.arange(-7, 7, 0.1)


```



```

phi_z = sigmoid(z)
plt.plot(z, phi_z)
plt.axvline(0.0, color='k')
plt.ylim(-0.1, 1.1)
plt.xlabel('z')
plt.ylabel('$\phi(z)$')
# y axis ticks and gridline
plt.yticks([0.0, 0.5, 1.0])
ax = plt.gca()
ax.yaxis.grid(True)
plt.show()

```




```

[48]: def cost_1(z):
        return - np.log(sigmoid(z))
    def cost_0(z):
        return - np.log(1 - sigmoid(z))
    z = np.arange(-10, 10, 0.1)
    phi_z = sigmoid(z)
    c1 = [cost_1(x) for x in z]
    plt.plot(phi_z, c1, label='J(w) if y=1')
    c0 = [cost_0(x) for x in z]

```

```
plt.plot(phi_z, c0, linestyle='--', label='J(w) if y=0')
plt.ylim(0.0, 5.1)
plt.xlim([0, 1])
plt.xlabel('$\phi(z)$')
plt.ylabel('J(w)')
plt.legend(loc='best')
plt.show()
```



0.2.1 Implementing Logistic Regression by modifying Adaline

```
[49]: class LogisticRegressionGD(object):
    """Logistic Regression Classifier using gradient descent.
    Parameters
    -----
    eta : float
    Learning rate (between 0.0 and 1.0)
    n_iter : int
    Passes over the training dataset.
    random_state : int
    Random number generator seed for random weight
```

```

initialization.
Attributes
-----
w_ : 1d-array
Weights after fitting.
cost_ : list
Sum-of-squares cost function value in each epoch.
"""
def __init__(self, eta=0.05, n_iter=100, random_state=1):
    self.eta = eta
    self.n_iter = n_iter
    self.random_state = random_state
def fit(self, X, y):
    """ Fit training data.
    Parameters
    -----
    X : {array-like}, shape = [n_samples, n_features]
    Training vectors, where n_samples is the number of
    samples and
    n_features is the number of features.
    y : array-like, shape = [n_samples]
    Target values.
    Returns
    -----
    self : object
    """
    rgen = np.random.RandomState(self.random_state)
    self.w_ = rgen.normal(loc=0.0, scale=0.01, size=1 + X.shape[1])
    self.cost_ = []
    for i in range(self.n_iter):
        net_input = self.net_input(X)
        output = self.activation(net_input)
        errors = (y - output)
        self.w_[1:] += self.eta * X.T.dot(errors)
        self.w_[0] += self.eta * errors.sum()
        # note that we compute the logistic `cost` now
        # instead of the sum of squared errors cost
        cost = (-y.dot(np.log(output)) -
                ((1 - y).dot(np.log(1 - output))))
        self.cost_.append(cost)
    return self
def net_input(self, X):
    """Calculate net input"""
    return np.dot(X, self.w_[1:]) + self.w_[0]
def activation(self, z):
    """Compute logistic sigmoid activation"""
    return 1. / (1. + np.exp(-np.clip(z, -250, 250)))

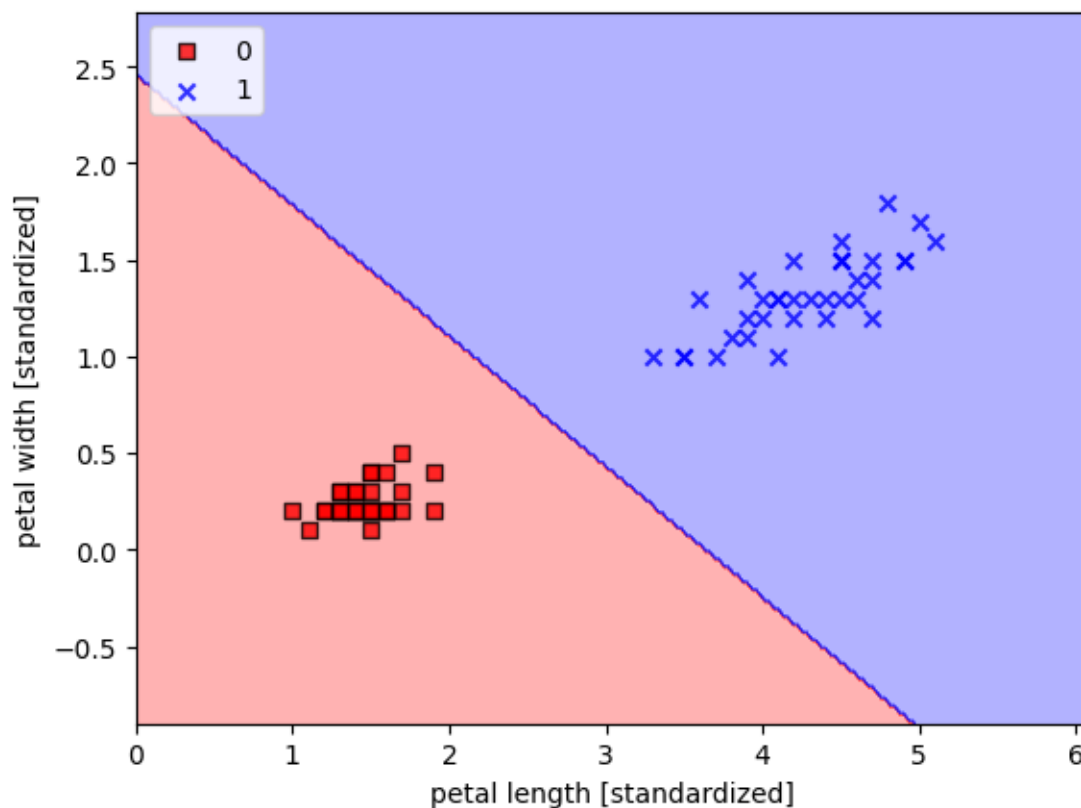
```

```
def predict(self, X):
    """Return class label after unit step"""
    return np.where(self.net_input(X) >= 0.0, 1, 0)
    # equivalent to:
    # return np.where(self.activation(self.net_input(X))
    # >= 0.5, 1, 0)
```

```
[50]: X_train_01_subset = X_train[(y_train == 0) | (y_train == 1)]
y_train_01_subset = y_train[(y_train == 0) | (y_train == 1)]
lrgd = LogisticRegressionGD(eta=0.05,n_iter=1000,random_state=1)
lrgd.fit(X_train_01_subset,y_train_01_subset)
plot_decision_regions(X=X_train_01_subset,y=y_train_01_subset,classifier=lrgd)
plt.xlabel('petal length [standardized]')
plt.ylabel('petal width [standardized]')
plt.legend(loc='upper left')
plt.show()
```

C:\Users\ankit19.gupta\AppData\Local\Temp\ipykernel_15304\1174551119.py:19:
UserWarning: You passed a edgecolor/edgecolors ('black') for an unfilled marker ('x'). Matplotlib is ignoring the edgecolor in favor of the facecolor. This behavior may change in the future.

```
plt.scatter(x=X[y == c1, 0], y=X[y == c1, 1],alpha=0.8,
c=colors[idx],marker=markers[idx], label=c1,edgecolor='black')
```




0.3 Training a logistic regression model with scikit-learn

```
[51]: from sklearn.linear_model import LogisticRegression
lr = LogisticRegression(C=100.0, random_state=1)
lr.fit(X_train_std, y_train)
plot_decision_regions(X_combined_std, y_combined, classifier=lr, test_idx=range(105, 150))
plt.xlabel('petal length [standardized]')
plt.ylabel('petal width [standardized]')
plt.legend(loc='upper left')
plt.show()
```

C:\Users\ankit19.gupta\AppData\Local\Temp\ipykernel_15304\1174551119.py:19:
UserWarning: You passed a edgecolor/edgecolors ('black') for an unfilled marker ('x'). Matplotlib is ignoring the edgecolor in favor of the facecolor. This behavior may change in the future.

```
plt.scatter(x=X[y == c1, 0], y=X[y == c1, 1], alpha=0.8,  
c=colors[idx], marker=markers[idx], label=c1, edgecolor='black')
```



```
[52]: lr.predict_proba(X_test_std[:3, :])

[52]: array([[9.37368464e-13, 3.91458193e-04, 9.99608542e-01],
           [9.93631074e-01, 6.36892585e-03, 1.20730798e-15],
           [9.98707332e-01, 1.29266792e-03, 1.82177043e-17]])

[53]: lr.predict_proba(X_test_std[:3, :]).sum(axis=1)

[53]: array([1., 1., 1.])

[54]: lr.predict_proba(X_test_std[:3, :]).argmax(axis=1)

[54]: array([2, 0, 0], dtype=int64)

[55]: lr.predict(X_test_std[:3, :])

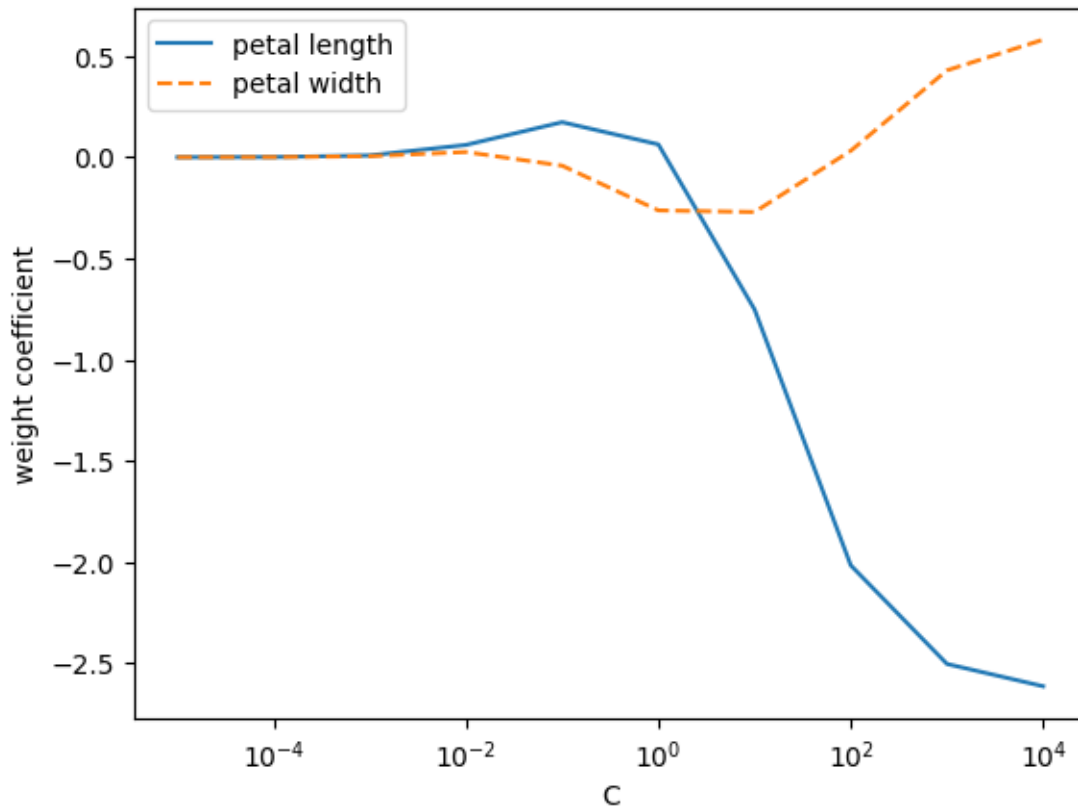
[55]: array([2, 0, 0])

[56]: lr.predict(X_test_std[0, :].reshape(1, -1))

[56]: array([2])
```

0.4 Tackling Overfitting and Regularization

```
[57]: weights, params = [], []
      for c in np.arange(-5, 5):
          lr = LogisticRegression(C=10.**c, random_state=1)
          lr.fit(X_train_std, y_train)
          weights.append(lr.coef_[1])
          params.append(10.**c)
      weights = np.array(weights)
      plt.plot(params, weights[:, 0], label='petal length')
      plt.plot(params, weights[:, 1], linestyle='--', label='petal width')
      plt.ylabel('weight coefficient')
      plt.xlabel('C')
      plt.legend(loc='upper left')
      plt.xscale('log')
      plt.show()
```




0.5 support vector machines

```
[58]: from sklearn.svm import SVC
svm = SVC(kernel='linear', C=1.0, random_state=1)
svm.fit(X_train_std, y_train)
plot_decision_regions(X_combined_std, y_combined, classifier=svm, test_idx=range(105, 150))
plt.xlabel('petal length [standardized]')
plt.ylabel('petal width [standardized]')
plt.legend(loc='upper left')
plt.show()
```

C:\Users\ankit19.gupta\AppData\Local\Temp\ipykernel_15304\1174551119.py:19:
UserWarning: You passed a edgecolor/edgecolors ('black') for an unfilled marker ('x'). Matplotlib is ignoring the edgecolor in favor of the facecolor. This behavior may change in the future.

```
plt.scatter(x=X[y == c1, 0], y=X[y == c1, 1], alpha=0.8,
c=colors[idx], marker=markers[idx], label=c1, edgecolor='black')
```



0.6 Alternative implementations in scikit-learn

```
[59]: # for large datasets

from sklearn.linear_model import SGDClassifier
ppn = SGDClassifier(loss='perceptron')
lr = SGDClassifier(loss='log')
svm = SGDClassifier(loss='hinge')
```


0.7 Solving nonlinear problems using a kernel SVM

0.7.1 Kernel methods for linearly inseparable data

```
[60]: import matplotlib.pyplot as plt
import numpy as np
np.random.seed(1)
X_xor = np.random.randn(200, 2)
y_xor = np.logical_xor(X_xor[:, 0] > 0, X_xor[:, 1] > 0)
y_xor = np.where(y_xor, 1, -1)
plt.scatter(X_xor[y_xor == 1, 0], X_xor[y_xor == 1, 1], c='b', s=10,
            ↪marker='x', label='1')
```




```
plt.scatter(X_xor[y_xor == -1, 0],X_xor[y_xor == -1,1],c='r',marker='s',label='-1')
plt.xlim([-3, 3])
plt.ylim([-3, 3])
plt.legend(loc='best')
plt.show()
```



```
[61]: svm = SVC(kernel='rbf', random_state=1, gamma=0.10, C=10.0)
svm.fit(X_xor, y_xor)
plot_decision_regions(X_xor, y_xor, classifier=svm)
plt.legend(loc='upper left')
plt.show()
```

C:\Users\ankit19.gupta\AppData\Local\Temp\ipykernel_15304\1174551119.py:19:
 UserWarning: You passed a edgecolor/edgecolors ('black') for an unfilled marker ('x'). Matplotlib is ignoring the edgecolor in favor of the facecolor. This behavior may change in the future.

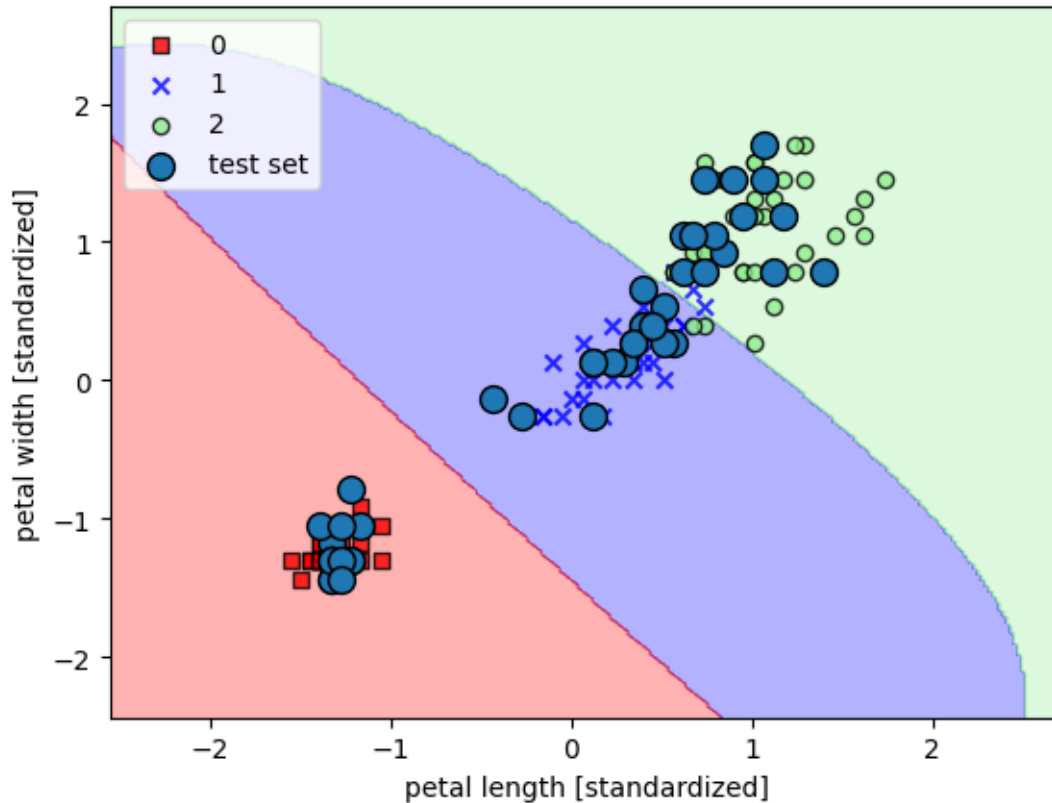
```
plt.scatter(x=X[y == c1, 0], y=X[y == c1, 1],alpha=0.8,
c=colors[idx],marker=markers[idx], label=c1,edgecolor='black')
```



```
[62]: svm = SVC(kernel='rbf', random_state=1, gamma=0.2, C=1.0)
svm.fit(X_train_std, y_train)
plot_decision_regions(X_combined_std, y_combined,
    classifier=svm, test_idx=range(105, 150))
plt.xlabel('petal length [standardized]')
plt.ylabel('petal width [standardized]')
plt.legend(loc='upper left')
plt.show()
```

C:\Users\ankit19.gupta\AppData\Local\Temp\ipykernel_15304\1174551119.py:19:
 UserWarning: You passed a edgecolor/edgecolors ('black') for an unfilled marker ('x'). Matplotlib is ignoring the edgecolor in favor of the facecolor. This behavior may change in the future.


```
plt.scatter(x=X[y == c1, 0], y=X[y == c1, 1], alpha=0.8,
c=colors[idx], marker=markers[idx], label=c1, edgecolor='black')
```



```
[63]: svm = SVC(kernel='rbf', random_state=1, gamma=100.0, C=1.0)
svm.fit(X_train_std, y_train)
plot_decision_regions(X_combined_std, y_combined,
    classifier=svm, test_idx=range(105, 150))
plt.xlabel('petal length [standardized]')
plt.ylabel('petal width [standardized]')
plt.legend(loc='upper left')
plt.show()
```

C:\Users\ankit19.gupta\AppData\Local\Temp\ipykernel_15304\1174551119.py:19:
 UserWarning: You passed a edgecolor/edgecolors ('black') for an unfilled marker ('x'). Matplotlib is ignoring the edgecolor in favor of the facecolor. This behavior may change in the future.

```
plt.scatter(x=X[y == c1, 0], y=X[y == c1, 1], alpha=0.8,
c=colors[idx], marker=markers[idx], label=c1, edgecolor='black')
```




0.8 Decision tree learning

```
[64]: import matplotlib.pyplot as plt
import numpy as np
def gini(p):
    return (p)*(1 - (p)) + (1 - p)*(1 - (1-p))
def entropy(p):
    return - p*np.log2(p) - (1 - p)*np.log2((1 - p))
def error(p):
    return 1 - np.max([p, 1 - p])
x = np.arange(0.0, 1.0, 0.01)
ent = [entropy(p) if p != 0 else None for p in x]
sc_ent = [e*0.5 if e else None for e in ent]
err = [error(i) for i in x]
fig = plt.figure()
ax = plt.subplot(111)
for i, lab, ls, c, in zip([ent, sc_ent, gini(x), err], ['Entropy', 'Entropy_
↳(scaled)', 'Gini Impurity', 'Misclassification Error'], ['-', '-', '---', '-.
↳'], ['black', 'lightgray', 'red', 'green', 'cyan']):
    line = ax.plot(x, i, label=lab, linestyle=ls, lw=2, color=c)
```

```

ax.legend(loc='upper center', bbox_to_anchor=(0.5, 1.15),ncol=5, fancybox=True,
        shadow=False)
ax.axhline(y=0.5, linewidth=1, color='k', linestyle='--')
ax.axhline(y=1.0, linewidth=1, color='k', linestyle='--')
plt.ylim([0, 1.1])
plt.xlabel('p(i=1)')
plt.ylabel('Impurity Index')
plt.show()

```



0.9 Building a decision tree

```


[65]: from sklearn.tree import DecisionTreeClassifier
tree = DecisionTreeClassifier(criterion='gini',max_depth=4,random_state=1)
tree.fit(X_train, y_train)
X_combined = np.vstack((X_train, X_test))
y_combined = np.hstack((y_train, y_test))
plot_decision_regions(X_combined,y_combined,classifier=tree,test_idx=range(105,
        150))
plt.xlabel('petal length [cm]')
plt.ylabel('petal width [cm]')
plt.legend(loc='upper left')

```

```
plt.show()
```

C:\Users\ankit19.gupta\AppData\Local\Temp\ipykernel_15304\1174551119.py:19:
UserWarning: You passed a edgecolor/edgecolors ('black') for an unfilled marker
(*'x'*). Matplotlib is ignoring the edgecolor in favor of the facecolor. This
behavior may change in the future.

```
plt.scatter(x=X[y == c1, 0], y=X[y == c1, 1],alpha=0.8,  
c=colors[idx],marker=markers[idx], label=c1,edgecolor='black')
```



```
[66]: # ! pip3 install pydotplus  
      # ! pip3 install graphviz  
      # ! pip3 install pyparsing
```

```
[68]: # first download executable from https://graphviz.org/download/ and then add_  
      ↪ into the path  
import os  
os.environ["PATH"] += os.pathsep + 'C:\Program Files (x86)\Graphviz-10.0.  
      ↪ 1-win64/bin/'
```

```
[69]: from pydotplus import graph_from_dot_data  
      from sklearn.tree import export_graphviz
```

```

dot_data = graphviz
    ↳ export_graphviz(tree,filled=True,rounded=True,class_names=['Setosa','Versicolor','Virginica',
    ↳ 'length','petal width'],out_file=None)
graph = graph_from_dot_data(dot_data)
graph.write_png('tree.png')

```

[69]: True

0.10 Combining multiple decision trees via random forests

```

[70]: from sklearn.ensemble import RandomForestClassifier
forest = RandomForestClassifier(criterion='gini',n_estimators=25,random_state=1,n_jobs=2)
forest.fit(X_train, y_train)
plot_decision_regions(X_combined, y_combined,classifier=forest,test_idx=range(105,150))
plt.xlabel('petal length')
plt.ylabel('petal width')
plt.legend(loc='upper left')
plt.show()


```

C:\Users\ankit19.gupta\AppData\Local\Temp\ipykernel_15304\1174551119.py:19:
UserWarning: You passed a edgecolor/edgecolors ('black') for an unfilled marker ('x'). Matplotlib is ignoring the edgecolor in favor of the facecolor. This behavior may change in the future.

```

plt.scatter(x=X[y == c1, 0], y=X[y == c1, 1],alpha=0.8,
c=colors[idx],marker=markers[idx], label=c1,edgecolor='black')

```



0.11 KNN - A Lazy Learning Algorithm

```
[71]: from sklearn.neighbors import KNeighborsClassifier
knn = KNeighborsClassifier(n_neighbors=5, p=2, metric='minkowski')
knn.fit(X_train_std, y_train)
plot_decision_regions(X_combined_std, y_combined, classifier=knn,
    ↪ test_idx=range(105, 150))
plt.xlabel('petal length [standardized]')
plt.ylabel('petal width [standardized]')
plt.legend(loc='upper left')
plt.show()
```

C:\Users\ankit19.gupta\AppData\Local\Temp\ipykernel_15304\1174551119.py:19:
UserWarning: You passed a edgecolor/edgecolors ('black') for an unfilled marker ('x'). Matplotlib is ignoring the edgecolor in favor of the facecolor. This behavior may change in the future.

```
plt.scatter(x=X[y == c1, 0], y=X[y == c1, 1], alpha=0.8,
c=colors[idx], marker=markers[idx], label=c1, edgecolor='black')
```