

Chapter_10_Predicting_Continuous_Target_Variables_with_Regression_A

March 19, 2024

0.1 Loading the Housing dataset into a data frame

```
[1]: import pandas as pd
df = pd.read_csv('https://raw.githubusercontent.com/rasbt/
python-machine-learning-book-2nd-edition/master/code/ch10/housing.data.
txt', header=None, sep='\s+')
df.columns = ['CRIM', 'ZN', 'INDUS', 'CHAS', 'NOX', 'RM', 'AGE', 'DIS', '
RAD', 'TAX', 'PTRATIO', 'B', 'LSTAT', 'MEDV']
df.head()
```

```
[1]:
```

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	\
0	0.00632	18.0	2.31	0	0.538	6.575	65.2	4.0900	1	296.0	
1	0.02731	0.0	7.07	0	0.469	6.421	78.9	4.9671	2	242.0	
2	0.02729	0.0	7.07	0	0.469	7.185	61.1	4.9671	2	242.0	
3	0.03237	0.0	2.18	0	0.458	6.998	45.8	6.0622	3	222.0	
4	0.06905	0.0	2.18	0	0.458	7.147	54.2	6.0622	3	222.0	

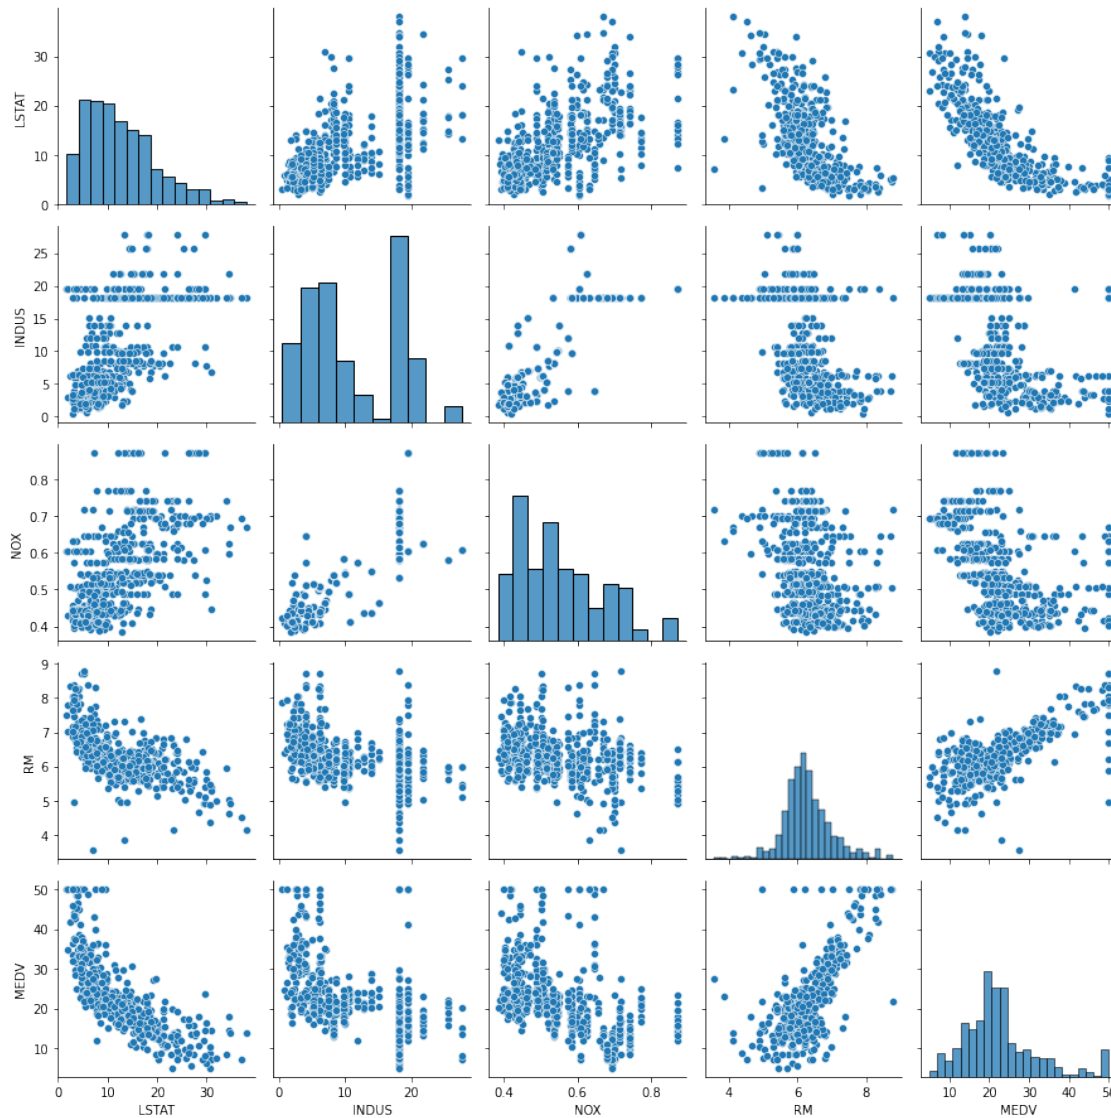
	PTRATIO	B	LSTAT	MEDV
0	15.3	396.90	4.98	24.0
1	17.8	396.90	9.14	21.6
2	17.8	392.83	4.03	34.7
3	18.7	394.63	2.94	33.4
4	18.7	396.90	5.33	36.2

0.2 Visualizing the important characteristics of a dataset

```
[4]: import matplotlib.pyplot as plt
import seaborn as sns
cols = ['LSTAT', 'INDUS', 'NOX', 'RM', 'MEDV']
sns.pairplot(df[cols], size=2.5)
plt.tight_layout()
plt.show()
```

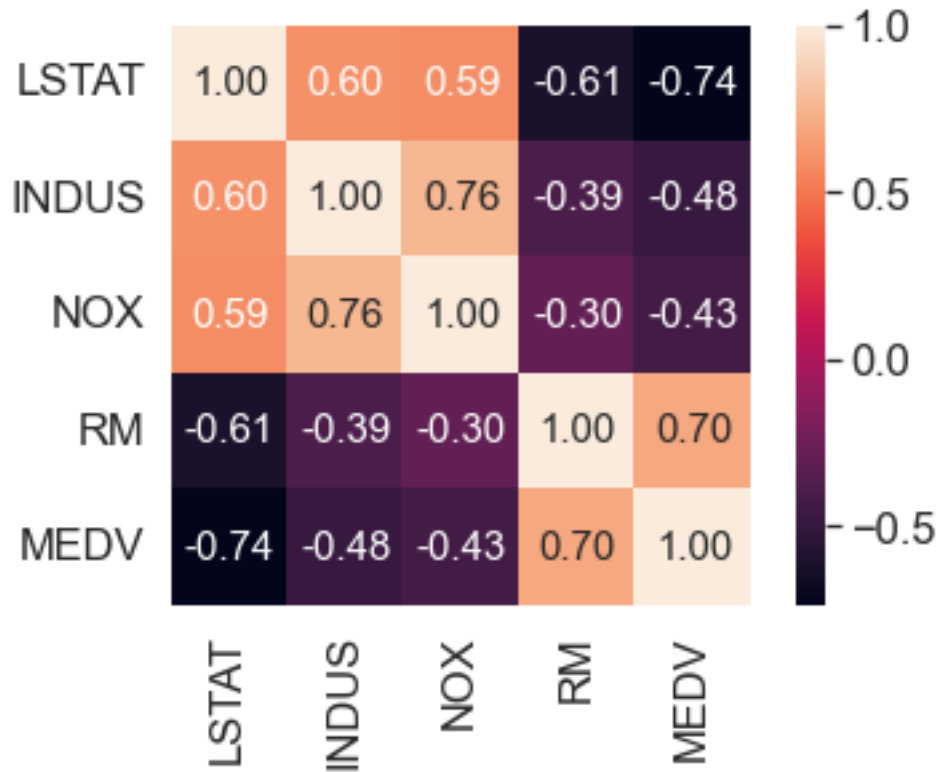
C:\Users\ankit19.gupta\OneDrive - Reliance Corporate IT Park Limited\Desktop\Sel
f_Projects\Python_Machine_Learning_Sebastian_Raschka\myenv\lib\site-
packages\seaborn\axisgrid.py:2076: UserWarning: The `size` parameter has been

renamed to `height`; please update your code.
 warnings.warn(msg, UserWarning)



0.3 Looking at relationships using a correlation matrix

```
[5]: import numpy as np
cm = np.corrcoef(df[cols].values.T)
sns.set(font_scale=1.5)
hm = sns.heatmap(cm, cbar=True, annot=True, square=True, fmt='.
    ↪2f', annot_kws={'size': 15}, yticklabels=cols, xticklabels=cols)
plt.show()
```



0.4 Implementing an ordinary least squares linear regression model

0.5 Solving regression for regression parameters with gradient descent


```
[6]: class LinearRegressionGD(object):
    def __init__(self, eta=0.001, n_iter=20):
        self.eta = eta
        self.n_iter = n_iter
    def fit(self, X, y):
        self.w_ = np.zeros(1 + X.shape[1])
        self.cost_ = []
        for i in range(self.n_iter):
            output = self.net_input(X)
            errors = (y - output)
            self.w_[1:] += self.eta * X.T.dot(errors)
            self.w_[0] += self.eta * errors.sum()
            cost = (errors**2).sum() / 2.0
            self.cost_.append(cost)
        return self
    def net_input(self, X):
        return np.dot(X, self.w_[1:]) + self.w_[0]
    def predict(self, X):
```

```
return self.net_input(X)
```

```
[7]: X = df[['RM']].values
y = df['MEDV'].values
from sklearn.preprocessing import StandardScaler
sc_x = StandardScaler()
sc_y = StandardScaler()
X_std = sc_x.fit_transform(X)
y_std = sc_y.fit_transform(y[:, np.newaxis]).flatten()
lr = LinearRegressionGD()
lr.fit(X_std, y_std)
```


```
[7]: <__main__.LinearRegressionGD at 0x2b5f4b2a240>
```

```
[8]: sns.reset_orig() # resets matplotlib style
plt.plot(range(1, lr.n_iter+1), lr.cost_)
plt.ylabel('SSE')
plt.xlabel('Epoch')
plt.show()
```



```
[9]: def lin_regplot(X, y, model):
    plt.scatter(X, y, c='steelblue', edgecolor='white', s=70)
    plt.plot(X, model.predict(X), color='black', lw=2)
    return None
```

```
[10]: lin_regplot(X_std, y_std, lr)
plt.xlabel('Average number of rooms [RM] (standardized)')
plt.ylabel('Price in $1000s [MEDV] (standardized)')
plt.show()
```



```
[15]: # num_rooms_std = sc_x.transform([5.0])
# price_std = lr.predict(num_rooms_std)
# print("Price in $1000s: %.3f" % sc_y.inverse_transform(price_std))
```

```
[16]: print('Slope: %.3f' % lr.w_[1])
print('Intercept: %.3f' % lr.w_[0])
```

Slope: 0.695
Intercept: -0.000

0.6 Estimating coefficients of the Regression model via scikit-learn


```
[18]: from sklearn.linear_model import LinearRegression
slr = LinearRegression()
slr.fit(X, y)
print('Slope: %.3f' % slr.coef_[0])
```

Slope: 9.102

```
[19]: print('Intercept: %.3f' % slr.intercept_)
```

Intercept: -34.671

```
[20]: lin_regplot(X, y, slr)
plt.xlabel('Average number of rooms [RM] (standardized)')
plt.ylabel('Price in $1000s [MEDV] (standardized)')
plt.show()
```



```
[21]: # adding a column vector of "ones"
Xb = np.hstack((np.ones((X.shape[0], 1)), X))
w = np.zeros(X.shape[1])
z = np.linalg.inv(np.dot(Xb.T, Xb))
w = np.dot(z, np.dot(Xb.T, y))
print('Slope: %.3f' % w[1])
print('Intercept: %.3f' % w[0])
```


Slope: 9.102
Intercept: -34.671

0.7 Fitting a robust regression model using RANSAC

```
[22]: from sklearn.linear_model import RANSACRegressor
ransac = RANSACRegressor(LinearRegression(), max_trials=100, min_samples=50, loss='absolute_loss', resid
ransac.fit(X, y)
```

```
[22]: RANSACRegressor(base_estimator=LinearRegression(), min_samples=50,  
                    random_state=0, residual_threshold=5.0)
```

```
[23]: inlier_mask = ransac.inlier_mask_  
outlier_mask = np.logical_not(inlier_mask)  
line_X = np.arange(3, 10, 1)  
line_y_ransac = ransac.predict(line_X[:, np.newaxis])  
plt.scatter(X[inlier_mask], y[inlier_mask], c='steelblue',  
            edgecolor='white', marker='o', label='Inliers')  
plt.scatter(X[outlier_mask], y[outlier_mask], c='limegreen',  
            edgecolor='white', marker='s', label='Outliers')  
plt.plot(line_X, line_y_ransac, color='black', lw=2)  
plt.xlabel('Average number of rooms [RM]')  
plt.ylabel('Price in $1000s [MEDV]')  
plt.legend(loc='upper left')  
plt.show()
```




```
[24]: print('Slope: %.3f' % ransac.estimator_.coef_[0])  
      print('Intercept: %.3f' % ransac.estimator_.intercept_)
```

Slope: 10.735
Intercept: -44.089

0.8 Evaluating the performance of linear regression models

```
[25]: from sklearn.model_selection import train_test_split
X = df.iloc[:, :-1].values
y = df['MEDV'].values
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3,
    random_state=0)
slr = LinearRegression()
slr.fit(X_train, y_train)
y_train_pred = slr.predict(X_train)
y_test_pred = slr.predict(X_test)
```

```
[26]: plt.scatter(y_train_pred, y_train_pred - y_train,
c='steelblue', marker='o', edgecolor='white',label='Training data')
plt.scatter(y_test_pred, y_test_pred - y_test,c='limegreen', marker='s',
    edgecolor='white',label='Test data')
plt.xlabel('Predicted values')
plt.ylabel('Residuals')
plt.legend(loc='upper left')
plt.hlines(y=0, xmin=-10, xmax=50, color='black', lw=2)
plt.xlim([-10, 50])
plt.show()
```




```
[27]: from sklearn.metrics import mean_squared_error
print('MSE train: %.3f, test: %.3f' % (mean_squared_error(y_train,
↪y_train_pred), mean_squared_error(y_test, y_test_pred)))
```

MSE train: 19.958, test: 27.196

```
[28]: from sklearn.metrics import r2_score
print('R^2 train: %.3f, test: %.3f' % (r2_score(y_train,
↪y_train_pred), r2_score(y_test, y_test_pred)))
```

R² train: 0.765, test: 0.673

0.9 Using regularized method for regression

```
[29]: from sklearn.linear_model import Ridge
ridge = Ridge(alpha=1.0)
```

```
[30]: from sklearn.linear_model import Lasso
lasso = Lasso(alpha=1.0)
```

```
[31]: from sklearn.linear_model import ElasticNet
elanet = ElasticNet(alpha=1.0, l1_ratio=0.5)
```

```
[32]: #For example, if we set the l1_ratio to 1.0, the ElasticNet regressor
#would be equal to LASSO regression
```

0.10 Turning a linear regression model into a curve-polynomial regression model


0.10.1 Adding polynomial terms using scikit-learn

```
[33]: from sklearn.preprocessing import PolynomialFeatures
X = np.array([ 258.0, 270.0, 294.0, 320.0, 342.0, 368.0, 396.0, 446.0, 480.0,
↪586.0])[:, np.newaxis]
y = np.array([ 236.4, 234.4, 252.8, 298.6, 314.2, 342.2, 360.8, 368.0, 391.2,
↪390.8])
lr = LinearRegression()
pr = LinearRegression()
quadratic = PolynomialFeatures(degree=2)
X_quad = quadratic.fit_transform(X)
```

```
[34]: lr.fit(X, y)
X_fit = np.arange(250, 600, 10)[:, np.newaxis]
y_lin_fit = lr.predict(X_fit)
```

```
[35]: pr.fit(X_quad, y)
y_quad_fit = pr.predict(quadratic.fit_transform(X_fit))
```

```
[36]: plt.scatter(X, y, label='training points')
plt.plot(X_fit, y_lin_fit, label='linear fit', linestyle='--')
plt.plot(X_fit, y_quad_fit, label='quadratic fit')
plt.legend(loc='upper left')
plt.show()
```



```
[37]: y_lin_pred = lr.predict(X)
y_quad_pred = pr.predict(X_quad)
print('Training MSE linear: %.3f, quadratic: %.3f' % (mean_squared_error(y, y_lin_pred), mean_squared_error(y, y_quad_pred)))
print('Training R^2 linear: %.3f, quadratic: %.3f' % (r2_score(y, y_lin_pred), r2_score(y, y_quad_pred)))
```

Training MSE linear: 569.780, quadratic: 61.330

Training R² linear: 0.832, quadratic: 0.982

0.11 Modeling nonlinear relationships in the Housing dataset


```
[38]: X = df[['LSTAT']].values
y = df['MEDV'].values
regr = LinearRegression()
# create quadratic features
quadratic = PolynomialFeatures(degree=2)
cubic = PolynomialFeatures(degree=3)
X_quad = quadratic.fit_transform(X)
X_cubic = cubic.fit_transform(X)
```

```


# fit features
X_fit = np.arange(X.min(), X.max(), 1)[: , np.newaxis]
regr = regr.fit(X, y)
y_lin_fit = regr.predict(X_fit)
linear_r2 = r2_score(y, regr.predict(X))
regr = regr.fit(X_quad, y)
y_quad_fit = regr.predict(quadratic.fit_transform(X_fit))
quadratic_r2 = r2_score(y, regr.predict(X_quad))
regr = regr.fit(X_cubic, y)
y_cubic_fit = regr.predict(cubic.fit_transform(X_fit))
cubic_r2 = r2_score(y, regr.predict(X_cubic))

# plot results
plt.scatter(X, y, label='training points', color='lightgray')
plt.plot(X_fit, y_lin_fit, label='linear (d=1), $R^2=%.2f$' % linear_r2, color='blue', lw=2, linestyle=':')
plt.plot(X_fit, y_quad_fit, label='quadratic (d=2), $R^2=%.2f$' % quadratic_r2, color='red', lw=2, linestyle='-')
plt.plot(X_fit, y_cubic_fit, label='cubic (d=3), $R^2=%.2f$' % cubic_r2, color='green', lw=2, linestyle='--')
plt.xlabel('% lower status of the population [LSTAT]')
plt.ylabel('Price in $1000s [MEDV]')
plt.legend(loc='upper right')
plt.show()

```



```
[39]: # transform features
X_log = np.log(X)
y_sqrt = np.sqrt(y)
# fit features
X_fit = np.arange(X_log.min()-1,X_log.max()+1, 1)[: , np.newaxis]
regr = regr.fit(X_log, y_sqrt)
y_lin_fit = regr.predict(X_fit)
linear_r2 = r2_score(y_sqrt, regr.predict(X_log))
# plot results
plt.scatter(X_log, y_sqrt,label='training points',color='lightgray')
plt.plot(X_fit, y_lin_fit,label='linear (d=1), $R^2=%.2f$' %
        linear_r2,color='blue',lw=2)
plt.xlabel('log(% lower status of the population [LSTAT])')
plt.ylabel('$\sqrt{\text{Price}}$ in $1000s $; [MEDV]$')
plt.legend(loc='lower left')
plt.show()
```




0.12 Dealing with nonlinear relationships using random forests

```
[40]: from sklearn.tree import DecisionTreeRegressor
X = df[['LSTAT']].values
y = df['MEDV'].values
tree = DecisionTreeRegressor(max_depth=3)
tree.fit(X, y)
```

```

sort_idx = X.flatten().argsort()
lin_regplot(X[sort_idx], y[sort_idx], tree)
plt.xlabel('% lower status of the population [LSTAT]')
plt.ylabel('Price in $1000s [MEDV]')
plt.show()

```



0.13 Random forest regression

```


[41]: X = df.iloc[:, :-1].values
y = df['MEDV'].values
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.
↳4, random_state=1)
from sklearn.ensemble import RandomForestRegressor
forest =
↳RandomForestRegressor(n_estimators=1000, criterion='mse', random_state=1, n_jobs=-1)
forest.fit(X_train, y_train)
y_train_pred = forest.predict(X_train)
y_test_pred = forest.predict(X_test)
print('MSE train: %.3f, test: %.3f' % (mean_squared_error(y_train,
↳y_train_pred), mean_squared_error(y_test, y_test_pred)))
print('R^2 train: %.3f, test: %.3f' % (
r2_score(y_train, y_train_pred),
r2_score(y_test, y_test_pred)))

```

MSE train: 1.644, test: 11.085

R² train: 0.979, test: 0.877

```
[42]: plt.scatter(y_train_pred,y_train_pred -  
    ↪y_train,c='steelblue',edgecolor='white',marker='o',s=35,alpha=0.  
    ↪9,label='Training data')  
plt.scatter(y_test_pred,y_test_pred -  
    ↪y_test,c='limegreen',edgecolor='white',marker='s',s=35,alpha=0.9,label='Test  
    ↪data')  
plt.xlabel('Predicted values')  
plt.ylabel('Residuals')  
plt.legend(loc='upper left')  
plt.hlines(y=0, xmin=-10, xmax=50, lw=2, color='black')  
plt.xlim([-10, 50])  
plt.show()
```



```
[ ]:
```