

Chapter_12_Implementing_a_Multilayer_Artificial_Neural_Network_from_

March 20, 2024

0.1 Modeling complex functions with Artificial Neural Networks

0.1.1 Classifying handwritten digits

```
[1]: import os
import struct
import numpy as np
```

```
[6]: def load_mnist(path, kind='train'):
    """Load MNIST data from `path`"""
    labels_path = os.path.join(path, '%s-labels.idx1-ubyte' % kind)
    images_path = os.path.join(path, '%s-images.idx3-ubyte' % kind)
    with open(labels_path, 'rb') as lbpath:
        magic, n = struct.unpack('>II', lbpath.read(8))
        labels = np.fromfile(lbpath, dtype=np.uint8)
    with open(images_path, 'rb') as imgpath:
        magic, num, rows, cols = struct.unpack(">IIII", imgpath.read(16))
        images = np.fromfile(imgpath, dtype=np.uint8).reshape(len(labels), 784)
        images = ((images / 255.) - .5) * 2
    return images, labels
```

```
[7]: X_train, y_train = load_mnist('', kind='train')
```

```
[8]: print('Rows: %d, columns: %d' % (X_train.shape[0], X_train.shape[1]))
```

Rows: 60000, columns: 784


```
[9]: X_test, y_test = load_mnist('', kind='t10k')
```

```
[10]: print('Rows: %d, columns: %d' % (X_test.shape[0], X_test.shape[1]))
```


Rows: 10000, columns: 784

To get an idea of how those images in MNIST look, let's visualize examples of the digits 0-9 after reshaping the 784-pixel vectors from our feature matrix into the original 28×28 image that we can plot via Matplotlib's `imshow` function:

```
[11]: import matplotlib.pyplot as plt
fig, ax = plt.subplots(nrows=2, ncols=5, sharex=True, sharey=True)
ax = ax.flatten()
for i in range(10):
    img = X_train[y_train == i][0].reshape(28, 28)
    ax[i].imshow(img, cmap='Greys')
ax[0].set_xticks([])
ax[0].set_yticks([])
plt.tight_layout()
plt.show()
```



```
[12]: fig, ax = plt.subplots(nrows=5, ncols=5, sharex=True, sharey=True)
ax = ax.flatten()
for i in range(25):
    img = X_train[y_train == 7][i].reshape(28, 28)
    ax[i].imshow(img, cmap='Greys')
ax[0].set_xticks([])
ax[0].set_yticks([])
plt.tight_layout()
plt.show()
```



```
[13]: import numpy as np
      np.savez_compressed('mnist_scaled.
      ↪npz', X_train=X_train, y_train=y_train, X_test=X_test, y_test=y_test)
```

```
[14]: mnist = np.load('mnist_scaled.npz')
```

```
[15]: mnist.files
```

```
[15]: ['X_train', 'y_train', 'X_test', 'y_test']
```

```
[16]: X_train = mnist['X_train']
```

```
[17]: X_train, y_train, X_test, y_test = [mnist[f] for f in mnist.files]
```

0.1.2 Implementing a multilayer perceptron

```
[18]: import numpy as np
      import sys
      class NeuralNetMLP(object):
          """ Feedforward neural network / Multi-layer perceptron
              classifier.
              Parameters
              -----
```

```

n_hidden : int (default: 30)
Number of hidden units.
l2 : float (default: 0.)
Lambda value for L2-regularization.
No regularization if l2=0. (default)
epochs : int (default: 100)
Number of passes over the training set.
eta : float (default: 0.001)
Learning rate.
shuffle : bool (default: True)
Shuffles training data every epoch
if True to prevent circles.
minibatch_size : int (default: 1)
Number of training samples per minibatch.
seed : int (default: None)
Random seed for initializing weights and shuffling.
Attributes
-----
eval_ : dict
Dictionary collecting the cost, training accuracy,
and validation accuracy for each epoch during training.
"""

def __init__(self, n_hidden=30, l2=0., epochs=100, eta=0.001, shuffle=True,
↪minibatch_size=1, seed=None):
    self.random = np.random.RandomState(seed)
    self.n_hidden = n_hidden
    self.l2 = l2
    self.epochs = epochs
    self.eta = eta
    self.shuffle = shuffle
    self.minibatch_size = minibatch_size
def _onehot(self, y, n_classes):
    """Encode labels into one-hot representation
    Parameters
    -----
    y : array, shape = [n_samples]
    Target values.
    Returns
    -----
    onehot : array, shape = (n_samples, n_labels)
    """
    onehot = np.zeros((n_classes, y.shape[0]))
    for idx, val in enumerate(y.astype(int)):
        onehot[val, idx] = 1.
    return onehot.T
def _sigmoid(self, z):
    """Compute logistic function (sigmoid)"""

```

```

        return 1. / (1. + np.exp(-np.clip(z, -250, 250)))
def _forward(self, X):
    """Compute forward propagation step"""
    # step 1: net input of hidden layer
    # [n_samples, n_features] dot [n_features, n_hidden]
    # -> [n_samples, n_hidden]
    z_h = np.dot(X, self.w_h) + self.b_h
    # step 2: activation of hidden layer
    a_h = self._sigmoid(z_h)
    # step 3: net input of output layer
    # [n_samples, n_hidden] dot [n_hidden, n_classlabels]
    # -> [n_samples, n_classlabels]
    z_out = np.dot(a_h, self.w_out) + self.b_out
    # step 4: activation output layer
    a_out = self._sigmoid(z_out)
    return z_h, a_h, z_out, a_out
def _compute_cost(self, y_enc, output):
    """Compute cost function.
    Parameters
    -----
    y_enc : array, shape = (n_samples, n_labels)
    one-hot encoded class labels.
    output : array, shape = [n_samples, n_output_units]
    Activation of the output layer (forward propagation)
    Returns
    -----
    cost : float
    Regularized cost
    """

    L2_term = (self.l2 * (np.sum(self.w_h ** 2.) + np.sum(self.w_out ** 2.)))
    term1 = -y_enc * (np.log(output))
    term2 = (1. - y_enc) * np.log(1. - output)
    cost = np.sum(term1 - term2) + L2_term
    return cost
def predict(self, X):
    """Predict class labels
    Parameters
    -----
    X : array, shape = [n_samples, n_features]
    Input layer with original features.
    Returns:
    -----
    y_pred : array, shape = [n_samples]
    Predicted class labels.
    """

    z_h, a_h, z_out, a_out = self._forward(X)
    y_pred = np.argmax(z_out, axis=1)

```

```

    return y_pred
def fit(self, X_train, y_train, X_valid, y_valid):
    """ Learn weights from training data.
    Parameters
    -----
    X_train : array, shape = [n_samples, n_features]
    ,PSOHPHQWLQJDOXOWLOD\HU$UWLÂFLDO1HXUDO1HWZRUNIURP6FUDWFK
    [ 400 ]
    Input layer with original features.
    y_train : array, shape = [n_samples]
    Target class labels.
    X_valid : array, shape = [n_samples, n_features]
    Sample features for validation during training
    y_valid : array, shape = [n_samples]
    Sample labels for validation during training
    Returns:
    -----
    self
    """
    n_output = np.unique(y_train).shape[0] # no. of class #labels
    n_features = X_train.shape[1]
    #####
    # Weight initialization
    #####
    # weights for input -> hidden
    self.b_h = np.zeros(self.n_hidden)
    self.w_h = self.random.normal(loc=0.0, scale=0.1,size=(n_features,self.
↪n_hidden))
    # weights for hidden -> output
    self.b_out = np.zeros(n_output)
    self.w_out = self.random.normal(loc=0.0, scale=0.1,size=(self.
↪n_hidden,n_output))
    epoch_strlen = len(str(self.epochs)) # for progr. format.
    self.eval_ = {'cost': [], 'train_acc': [], 'valid_acc': \
    []}
    y_train_enc = self._onehot(y_train, n_output)
    # iterate over training epochs
    for i in range(self.epochs):
        # iterate over minibatches
        indices = np.arange(X_train.shape[0])
        if self.shuffle:
            self.random.shuffle(indices)
        for start_idx in range(0, indices.shape[0] -self.minibatch_size +1,
↪self.minibatch_size):
            batch_idx = indices[start_idx:start_idx +self.minibatch_size]
            # forward propagation
            z_h, a_h, z_out, a_out = self._forward(X_train[batch_idx])

```

```

#####
# Backpropagation
#####
# [n_samples, n_classlabels]
sigma_out = a_out - y_train_enc[batch_idx]
# [n_samples, n_hidden]
sigmoid_derivative_h = a_h * (1. - a_h)
# [n_samples, n_classlabels] dot [n_classlabels,
# n_hidden]
# -> [n_samples, n_hidden]
sigma_h = (np.dot(sigma_out, self.w_out.T)
↳*sigmoid_derivative_h)
# [n_features, n_samples] dot [n_samples,
# n_hidden]
# -> [n_features, n_hidden]
grad_w_h = np.dot(X_train[batch_idx].T, sigma_h)
grad_b_h = np.sum(sigma_h, axis=0)
# [n_hidden, n_samples] dot [n_samples,
# n_classlabels]
# -> [n_hidden, n_classlabels]
grad_w_out = np.dot(a_h.T, sigma_out)
grad_b_out = np.sum(sigma_out, axis=0)
# Regularization and weight updates
delta_w_h = (grad_w_h + self.l2*self.w_h)
delta_b_h = grad_b_h # bias is not regularized
self.w_h -= self.eta * delta_w_h
self.b_h -= self.eta * delta_b_h
delta_w_out = (grad_w_out + self.l2*self.w_out)
delta_b_out = grad_b_out # bias is not regularized
self.w_out -= self.eta * delta_w_out
self.b_out -= self.eta * delta_b_out
#####
# Evaluation
#####
# Evaluation after each epoch during training
z_h, a_h, z_out, a_out = self._forward(X_train)

cost = self._compute_cost(y_enc=y_train_enc,output=a_out)
y_train_pred = self.predict(X_train)
y_valid_pred = self.predict(X_valid)
train_acc = ((np.sum(y_train ==y_train_pred)).astype(np.float) /
↳X_train.shape[0])
valid_acc = ((np.sum(y_valid ==y_valid_pred)).astype(np.float) /
↳X_valid.shape[0])
sys.stderr.write('\r%0*d/%d | Cost: %.2f | Train/Valid Acc.: %.2f%%/
↳%.2f%%'%(epoch_strlen, i+1, self.epochs,cost,train_acc*100, valid_acc*100))
sys.stderr.flush()

```

```

        self.eval_['cost'].append(cost)
        self.eval_['train_acc'].append(train_acc)
        self.eval_['valid_acc'].append(valid_acc)
    return self

```

```

[19]: nn = NeuralNetMLP(n_hidden=100,l2=0.01,epochs=200,eta=0.
      ↪0005,minibatch_size=100,shuffle=True,seed=1)

```

```

[20]: nn.fit(X_train=X_train[:55000],y_train=y_train[:55000],X_valid=X_train[55000:
      ↪],y_valid=y_train[55000:])

```

200/200 | Cost: 5065.78 | Train/Valid Acc.: 99.28%/97.98%

```


[20]: <__main__.NeuralNetMLP at 0x1e701dcff60>

```

```

[21]: import matplotlib.pyplot as plt
      plt.plot(range(nn.epochs), nn.eval_['cost'])
      plt.ylabel('Cost')
      plt.xlabel('Epochs')
      plt.show()

```




```

[22]: plt.plot(range(nn.epochs), nn.eval_['train_acc'],label='training')
      plt.plot(range(nn.epochs), nn.eval_['valid_acc'],label='validation',
      ↪linestyle='--')
      plt.ylabel('Accuracy')

```







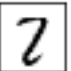




















```
plt.xlabel('Epochs')
plt.legend()
plt.show()
```



```
[23]: y_test_pred = nn.predict(X_test)
acc = (np.sum(y_test == y_test_pred).astype(np.float) / X_test.shape[0])
print('Training accuracy: %.2f%%' % (acc * 100))
```

Training accuracy: 97.54%

```
[29]: miscl_img = X_test[y_test != y_test_pred][:25]
correct_lab = y_test[y_test != y_test_pred][:25]
miscl_lab= y_test_pred[y_test != y_test_pred][:25]
fig, ax = plt.subplots(nrows=5,ncols=5,sharex=True,sharey=True,)
ax = ax.flatten()
for i in range(25):
    img = miscl_img[i].reshape(28, 28)
    ax[i].imshow(img,cmap='Greys',interpolation='nearest')
    ax[i].set_title('%d t: %d p: %d' % (i+1, correct_lab[i], miscl_lab[i]))
ax[0].set_xticks([])
ax[0].set_yticks([])
plt.tight_layout()
plt.show()
```

- 1) t: 5 p: 6 2) t: 4 p: 9 3) t: 4 p: 2 4) t: 6 p: 0 5) t: 2 p: 7
    
- 6) t: 5 p: 3 7) t: 3 p: 7 8) t: 6 p: 0 9) t: 3 p: 5 10) t: 8 p: 0
    
- 11) t: 7 p: 1 12) t: 3 p: 7 13) t: 1 p: 8 14) t: 2 p: 6 15) t: 2 p: 8
    
- 16) t: 7 p: 3 17) t: 8 p: 4 18) t: 5 p: 8 19) t: 4 p: 9 20) t: 9 p: 7
    
- 21) t: 2 p: 7 22) t: 3 p: 5 23) t: 8 p: 9 24) t: 5 p: 4 25) t: 1 p: 2
    

[]: