

Chapter_15_Classifying_Images_with_Deep_Convolutional_Neural_Netwo

March 20, 2024

```
[1]: import numpy as np
def conv1d(x, w, p=0, s=1):
    w_rot = np.array(w[::-1])
    x_padded = np.array(x)
    if p > 0:
        zero_pad = np.zeros(shape=p)
        x_padded = np.concatenate([zero_pad, x_padded, zero_pad])
    res = []
    for i in range(0, int(len(x)/s), s):
        res.append(np.sum(x_padded[i:i+w_rot.shape[0]] * w_rot))
    return np.array(res)
```

```
[2]: ## Testing:
x = [1, 3, 2, 4, 5, 6, 1, 3]
w = [1, 0, 3, 1, 2]
```

```
[3]: print('Conv1d Implementation:', conv1d(x, w, p=2, s=1))
print('Numpy Results:', np.convolve(x, w, mode='same'))
```

Conv1d Implementation: [5. 14. 16. 26. 24. 34. 19. 22.]

Numpy Results: [5 14 16 26 24 34 19 22]

```
[4]: import numpy as np
import scipy.signal
def conv2d(X, W, p=(0, 0), s=(1, 1)):
    W_rot = np.array(W)[::-1,::-1]
    X_orig = np.array(X)
    n1 = X_orig.shape[0] + 2*p[0]
    n2 = X_orig.shape[1] + 2*p[1]
    X_padded = np.zeros(shape=(n1, n2))
    X_padded[p[0]:p[0]+X_orig.shape[0], p[1]:p[1]+X_orig.shape[1]] = X_orig
    res = []
    for i in range(0, int((X_padded.shape[0] - W_rot.shape[0])/s[0])+1, s[0]):
        res.append([])
        for j in range(0, int((X_padded.shape[1] - W_rot.shape[1])/s[1])+1, s[1]):
```

```

        X_sub = X_padded[i:i+W_rot.shape[0],j:j+W_rot.shape[1]]
        res[-1].append(np.sum(X_sub * W_rot))
    return(np.array(res))

```

```

[5]: X = [[1, 3, 2, 4], [5, 6, 1, 3], [1, 2, 0, 2], [3, 4, 3, 2]]
     W = [[1, 0, 3], [1, 2, 1], [0, 1, 1]]

```

```

[6]: print('Conv2d Implementation:\n',conv2d(X, W, p=(1, 1), s=(1, 1)))

```

```

Conv2d Implementation:
[[11. 25. 32. 13.]
 [19. 25. 24. 13.]
 [13. 28. 25. 17.]
 [11. 17. 14.  9.]]

```

```

[7]: print('SciPy Results:\n',scipy.signal.convolve2d(X, W, mode='same'))

```

```

SciPy Results:
[[11 25 32 13]
 [19 25 24 13]
 [13 28 25 17]
 [11 17 14  9]]

```

```

[8]: # ! pip install scipy==1.1.0 --user

```

```

[9]: import scipy.misc
     import imageio
     #img = imageio.imread('./example-image.png',pilemode='RGB')
     img = scipy.misc.imread('./example-image.png',mode='RGB')
     print('Image shape:', img.shape)
     print('Number of channels:', img.shape[2])
     print('Image data type:', img.dtype)
     print(img[100:102, 100:102, :])

```

```

Image shape: (252, 221, 3)
Number of channels: 3
Image data type: uint8
[[[179 134 110]
  [182 136 112]]

 [[180 135 111]
  [182 137 113]]]

```

```

C:\Users\ankit19.gupta\OneDrive - Reliance Corporate IT Park Limited\Desktop\Sel
f_Projects\Python_Machine_Learning_Sebastian_Raschka\myenv\lib\site-
packages\ipykernel_launcher.py:4: DeprecationWarning:      `imread` is
deprecated!

```

```

    `imread` is deprecated in SciPy 1.0.0, and will be removed in 1.2.0.
    Use ``imageio.imread`` instead.

```

after removing the cwd from sys.path.

```
[11]: import os
import struct
import numpy as np
def load_mnist(path, kind='train'):
    """Load MNIST data from `path`"""
    labels_path = os.path.join(path, '%s-labels.idx1-ubyte' % kind)
    images_path = os.path.join(path, '%s-images.idx3-ubyte' % kind)
    with open(labels_path, 'rb') as lbpath:
        magic, n = struct.unpack('>II', lbpath.read(8))
        labels = np.fromfile(lbpath, dtype=np.uint8)
    with open(images_path, 'rb') as imgpath:
        magic, num, rows, cols = struct.unpack(">IIII", imgpath.read(16))
        images = np.fromfile(imgpath, dtype=np.uint8).reshape(len(labels), 784)
        images = ((images / 255.) - .5) * 2
    return images, labels
```

```
[12]: ##### Loading the data
X_data, y_data = load_mnist('./mnist/', kind='train')
print('Rows: {}, Columns: {}'.format(X_data.shape[0], X_data.shape[1]))
X_test, y_test = load_mnist('./mnist/', kind='t10k')
print('Rows: {}, Columns: {}'.format(X_test.shape[0], X_test.shape[1]))
X_train, y_train = X_data[:50000, :], y_data[:50000]
X_valid, y_valid = X_data[50000: :, :], y_data[50000:]
print('Training: ', X_train.shape, y_train.shape)
print('Validation: ', X_valid.shape, y_valid.shape)
print('Test Set: ', X_test.shape, y_test.shape)
```

```
Rows: 60000, Columns: 784
Rows: 10000, Columns: 784
Training: (50000, 784) (50000,)
Validation: (10000, 784) (10000,)
Test Set: (10000, 784) (10000,)
```

```
[13]: def batch_generator(X, y, batch_size=64, shuffle=False, random_seed=None):
    idx = np.arange(y.shape[0])
    if shuffle:
        rng = np.random.RandomState(random_seed)
        rng.shuffle(idx)
        X = X[idx]
        y = y[idx]
    for i in range(0, X.shape[0], batch_size):
        yield (X[i:i+batch_size, :], y[i:i+batch_size])
```

```
[14]: mean_vals = np.mean(X_train, axis=0)
std_val = np.std(X_train)
X_train_centered = (X_train - mean_vals)/std_val
```

```
X_valid_centered = (X_valid - mean_vals)/std_val
X_test_centered = (X_test - mean_vals)/std_val
```

```
[15]: import tensorflow as tf
import numpy as np
def conv_layer(input_tensor, name, kernel_size,
    ↪ n_output_channels, padding_mode='SAME', strides=(1, 1, 1, 1)):
    with tf.variable_scope(name):
        ## get n_input_channels:
        ## input tensor shape:
        ## [batch x width x height x channels_in]
        input_shape = input_tensor.get_shape().as_list()
        n_input_channels = input_shape[-1]
        weights_shape = list(kernel_size) + [n_input_channels,
    ↪ n_output_channels]
        weights = tf.get_variable(name='_weights', shape=weights_shape)
        print(weights)
        biases = tf.get_variable(name='_biases', initializer=tf.
    ↪ zeros(shape=[n_output_channels]))
        print(biases)
        conv = tf.nn.
    ↪ conv2d(input=input_tensor, filter=weights, strides=strides, padding=padding_mode)
        print(conv)
        conv = tf.nn.bias_add(conv, biases, name='net_pre-activation')
        print(conv)
        conv = tf.nn.relu(conv, name='activation')
        print(conv)

    return conv
```

C:\Users\ankit19.gupta\OneDrive - Reliance Corporate IT Park Limited\Desktop\Self_Projects\Python_Machine_Learning_Sebastian_Raschka\myenv\lib\site-packages\tensorflow\python\framework\dtypes.py:458: FutureWarning: Passing (type, 1) or '1type' as a synonym of type is deprecated; in a future version of numpy, it will be understood as (type, (1,)) / '(1,)type'.

```
_np_qint8 = np.dtype(["qint8", np.int8, 1])
```

C:\Users\ankit19.gupta\OneDrive - Reliance Corporate IT Park Limited\Desktop\Self_Projects\Python_Machine_Learning_Sebastian_Raschka\myenv\lib\site-packages\tensorflow\python\framework\dtypes.py:459: FutureWarning: Passing (type, 1) or '1type' as a synonym of type is deprecated; in a future version of numpy, it will be understood as (type, (1,)) / '(1,)type'.

```
_np_quint8 = np.dtype(["quint8", np.uint8, 1])
```

C:\Users\ankit19.gupta\OneDrive - Reliance Corporate IT Park Limited\Desktop\Self_Projects\Python_Machine_Learning_Sebastian_Raschka\myenv\lib\site-packages\tensorflow\python\framework\dtypes.py:460: FutureWarning: Passing (type, 1) or '1type' as a synonym of type is deprecated; in a future version of numpy, it will be understood as (type, (1,)) / '(1,)type'.

```

_np_qint16 = np.dtype([("qint16", np.int16, 1)])
C:\Users\ankit19.gupta\OneDrive - Reliance Corporate IT Park Limited\Desktop\Self_Projects\Python_Machine_Learning_Sebastian_Raschka\myenv\lib\site-packages\tensorflow\python\framework\dtypes.py:461: FutureWarning: Passing (type, 1) or '1type' as a synonym of type is deprecated; in a future version of numpy, it will be understood as (type, (1,)) / '(1,)type'.
_np_quint16 = np.dtype([("quint16", np.uint16, 1)])
C:\Users\ankit19.gupta\OneDrive - Reliance Corporate IT Park Limited\Desktop\Self_Projects\Python_Machine_Learning_Sebastian_Raschka\myenv\lib\site-packages\tensorflow\python\framework\dtypes.py:462: FutureWarning: Passing (type, 1) or '1type' as a synonym of type is deprecated; in a future version of numpy, it will be understood as (type, (1,)) / '(1,)type'.
_np_qint32 = np.dtype([("qint32", np.int32, 1)])
C:\Users\ankit19.gupta\OneDrive - Reliance Corporate IT Park Limited\Desktop\Self_Projects\Python_Machine_Learning_Sebastian_Raschka\myenv\lib\site-packages\tensorflow\python\framework\dtypes.py:465: FutureWarning: Passing (type, 1) or '1type' as a synonym of type is deprecated; in a future version of numpy, it will be understood as (type, (1,)) / '(1,)type'.
_np_resource = np.dtype([("resource", np.ubyte, 1)])

```

```

[16]: g = tf.Graph()
      with g.as_default():
          x = tf.placeholder(tf.float32, shape=[None, 28, 28, 1])
          conv_layer(x, name='convtest', kernel_size=(3, 3), n_output_channels=32)
      del g, x

```

```

<tf.Variable 'convtest/_weights:0' shape=(3, 3, 1, 32) dtype=float32_ref>
<tf.Variable 'convtest/_biases:0' shape=(32,) dtype=float32_ref>
Tensor("convtest/Conv2D:0", shape=(?, 28, 28, 32), dtype=float32)
Tensor("convtest/net_pre-activation:0", shape=(?, 28, 28, 32), dtype=float32)
Tensor("convtest/activation:0", shape=(?, 28, 28, 32), dtype=float32)

```

```

[17]: def fc_layer(input_tensor, name, n_output_units, activation_fn=None):
      with tf.variable_scope(name):
          input_shape = input_tensor.get_shape().as_list()[1:]
          n_input_units = np.prod(input_shape)
          if len(input_shape) > 1:
              input_tensor = tf.reshape(input_tensor, shape=(-1, n_input_units))
          weights_shape = [n_input_units, n_output_units]
          weights = tf.get_variable(name='_weights', shape=weights_shape)
          print(weights)
          biases = tf.get_variable(name='_biases', initializer=tf.
↳ zeros(shape=[n_output_units]))
          print(biases)
          layer = tf.matmul(input_tensor, weights)
          print(layer)
          layer = tf.nn.bias_add(layer, biases, name='net_pre-activation')
          print(layer)

```

```

    if activation_fn is None:
        return layer

    layer = activation_fn(layer, name='activation')
    print(layer)
    return layer

```

```

[18]: g = tf.Graph()
      with g.as_default():
          x = tf.placeholder(tf.float32, shape=[None, 28, 28, 1])
          fc_layer(x, name='fctest', n_output_units=32, activation_fn=tf.nn.relu)
      del g, x

```

```

<tf.Variable 'fctest/_weights:0' shape=(784, 32) dtype=float32_ref>
<tf.Variable 'fctest/_biases:0' shape=(32,) dtype=float32_ref>
Tensor("fctest/MatMul:0", shape=(?, 32), dtype=float32)
Tensor("fctest/net_pre-activation:0", shape=(?, 32), dtype=float32)
Tensor("fctest/activation:0", shape=(?, 32), dtype=float32)

```

```

[19]: def build_cnn():
      ## Placeholders for X and y:
      tf_x = tf.placeholder(tf.float32, shape=[None, 784], name='tf_x')
      tf_y = tf.placeholder(tf.int32, shape=[None], name='tf_y')
      # reshape x to a 4D tensor:
      # [batchsize, width, height, 1]
      tf_x_image = tf.reshape(tf_x, shape=[-1, 28, 28, 1], name='tf_x_reshaped')
      ## One-hot encoding:
      tf_y_onehot = tf.one_hot(indices=tf_y, depth=10, dtype=tf.
↪ float32, name='tf_y_onehot')
      ## 1st layer: Conv_1
      print('\nBuilding 1st layer:')
      h1 = conv_layer(tf_x_image, name='conv_1', kernel_size=(5,
↪ 5), padding_mode='VALID', n_output_channels=32)
      ## MaxPooling
      h1_pool = tf.nn.max_pool(h1, ksize=[1, 2, 2, 1], strides=[1, 2, 2,
↪ 1], padding='SAME')
      ## 2nd layer: Conv_2
      print('\nBuilding 2nd layer:')
      h2 = conv_layer(h1_pool, name='conv_2', kernel_size=(5,
↪ 5), padding_mode='VALID', n_output_channels=64)
      ## MaxPooling
      h2_pool = tf.nn.max_pool(h2, ksize=[1, 2, 2, 1], strides=[1, 2, 2,
↪ 1], padding='SAME')
      ## 3rd layer: Fully Connected
      print('\nBuilding 3rd layer:')
      h3 = fc_layer(h2_pool, name='fc_3', n_output_units=1024, activation_fn=tf.nn.
↪ relu)

```

```

## Dropout
keep_prob = tf.placeholder(tf.float32, name='fc_keep_prob')
h3_drop = tf.nn.dropout(h3, keep_prob,name='dropout_layer')
## 4th layer: Fully Connected (linear activation)
print('\nBuilding 4th layer:')
h4 = fc_layer(h3_drop, name='fc_4',n_output_units=10,activation_fn=None)
## Prediction
predictions = {
    'probabilities': tf.nn.softmax(h4, name='probabilities'),
    'labels': tf.cast(tf.argmax(h4, axis=1), tf.int32,
name='labels')
}

## Visualize the graph with TensorBoard:
## Loss Function and Optimization
cross_entropy_loss = tf.reduce_mean(tf.nn.
↪softmax_cross_entropy_with_logits(logits=h4,
↪labels=tf_y_onehot),name='cross_entropy_loss')
## Optimizer:
optimizer = tf.train.AdamOptimizer(learning_rate)
optimizer = optimizer.minimize(cross_entropy_loss,name='train_op')
## Computing the prediction accuracy
correct_predictions = tf.equal(predictions['labels'],tf_y,
↪name='correct_preds')
accuracy = tf.reduce_mean(tf.cast(correct_predictions, tf.
↪float32),name='accuracy')

```

```

[20]: def save(saver, sess, epoch, path='./model/'):
    if not os.path.isdir(path):
        os.makedirs(path)
    print('Saving model in %s' % path)
    saver.save(sess, os.path.join(path, 'cnn-model.ckpt'),global_step=epoch)
def load(saver, sess, path, epoch):
    print('Loading model from %s' % path)
    saver.restore(sess, os.path.join(path, 'cnn-model.ckpt-%d' % epoch))
def train(sess, training_set, validation_set=None,initialize=True, epochs=20,
↪shuffle=True,dropout=0.5, random_seed=None):
    X_data = np.array(training_set[0])
    y_data = np.array(training_set[1])
    training_loss = []
    ## initialize variables
    if initialize:
        sess.run(tf.global_variables_initializer())
    np.random.seed(random_seed) # for shuffling in batch_generator
    for epoch in range(1, epochs+1):
        batch_gen = batch_generator(X_data, y_data,shuffle=shuffle)
        avg_loss = 0.0

```

```

        for i,(batch_x,batch_y) in enumerate(batch_gen):
            feed = {'tf_x:0': batch_x,'tf_y:0': batch_y,'fc_keep_prob:0':
↳dropout}
            loss, _ = sess.run(['cross_entropy_loss:0',
↳'train_op'],feed_dict=feed)
            avg_loss += loss
            training_loss.append(avg_loss / (i+1))
            print('Epoch %02d Training Avg. Loss: %7.3f' % (epoch, avg_loss), end='
↳')

            if validation_set is not None:
                feed = {'tf_x:0': validation_set[0],'tf_y:0':
↳validation_set[1],'fc_keep_prob:0': 1.0}
                valid_acc = sess.run('accuracy:0', feed_dict=feed)
                print(' Validation Acc: %7.3f' % valid_acc)
            else:
                print()

def predict(sess, X_test, return_proba=False):
    feed = {'tf_x:0': X_test,'fc_keep_prob:0': 1.0}
    if return_proba:
        return sess.run('probabilities:0', feed_dict=feed)
    else:
        return sess.run('labels:0', feed_dict=feed)

```

```

[21]: ## Define hyperparameters
learning_rate = 1e-4
random_seed = 123
## create a graph
g = tf.Graph()
with g.as_default():
    tf.set_random_seed(random_seed)
    ## build the graph
    build_cnn()
    ## saver:
    saver = tf.train.Saver()

```

Building 1st layer:

```

<tf.Variable 'conv_1/_weights:0' shape=(5, 5, 1, 32) dtype=float32_ref>
<tf.Variable 'conv_1/_biases:0' shape=(32,) dtype=float32_ref>
Tensor("conv_1/Conv2D:0", shape=(?, 24, 24, 32), dtype=float32)
Tensor("conv_1/net_pre-activation:0", shape=(?, 24, 24, 32), dtype=float32)
Tensor("conv_1/activation:0", shape=(?, 24, 24, 32), dtype=float32)

```

Building 2nd layer:

```

<tf.Variable 'conv_2/_weights:0' shape=(5, 5, 32, 64) dtype=float32_ref>
<tf.Variable 'conv_2/_biases:0' shape=(64,) dtype=float32_ref>
Tensor("conv_2/Conv2D:0", shape=(?, 8, 8, 64), dtype=float32)

```



```
Tensor("conv_2/net_pre-activation:0", shape=(?, 8, 8, 64), dtype=float32)
Tensor("conv_2/activation:0", shape=(?, 8, 8, 64), dtype=float32)
```

Building 3rd layer:

```
<tf.Variable 'fc_3/_weights:0' shape=(1024, 1024) dtype=float32_ref>
<tf.Variable 'fc_3/_biases:0' shape=(1024,) dtype=float32_ref>
Tensor("fc_3/MatMul:0", shape=(?, 1024), dtype=float32)
Tensor("fc_3/net_pre-activation:0", shape=(?, 1024), dtype=float32)
Tensor("fc_3/activation:0", shape=(?, 1024), dtype=float32)
```

Building 4th layer:

```
<tf.Variable 'fc_4/_weights:0' shape=(1024, 10) dtype=float32_ref>
<tf.Variable 'fc_4/_biases:0' shape=(10,) dtype=float32_ref>
Tensor("fc_4/MatMul:0", shape=(?, 10), dtype=float32)
Tensor("fc_4/net_pre-activation:0", shape=(?, 10), dtype=float32)
```

```
[22]: ## create a TF session
      ## and train the CNN model
      with tf.Session(graph=g) as sess:
          train(sess, training_set=(X_train_centered,
      ↪ y_train), validation_set=(X_valid_centered,
      ↪ y_valid), initialize=True, random_seed=123)
          save(saver, sess, epoch=20)
```

```
Epoch 01 Training Avg. Loss: 273.649 Validation Acc: 0.974
Epoch 02 Training Avg. Loss: 76.096 Validation Acc: 0.983
Epoch 03 Training Avg. Loss: 50.239 Validation Acc: 0.986
Epoch 04 Training Avg. Loss: 39.219 Validation Acc: 0.986
Epoch 05 Training Avg. Loss: 32.226 Validation Acc: 0.989
Epoch 06 Training Avg. Loss: 27.557 Validation Acc: 0.989
Epoch 07 Training Avg. Loss: 23.365 Validation Acc: 0.990
Epoch 08 Training Avg. Loss: 19.373 Validation Acc: 0.990
Epoch 09 Training Avg. Loss: 17.170 Validation Acc: 0.991
Epoch 10 Training Avg. Loss: 14.973 Validation Acc: 0.992
Epoch 11 Training Avg. Loss: 12.827 Validation Acc: 0.990
Epoch 12 Training Avg. Loss: 11.740 Validation Acc: 0.992
Epoch 13 Training Avg. Loss: 10.098 Validation Acc: 0.992
Epoch 14 Training Avg. Loss: 8.866 Validation Acc: 0.992
Epoch 15 Training Avg. Loss: 7.439 Validation Acc: 0.990
Epoch 16 Training Avg. Loss: 6.786 Validation Acc: 0.990
Epoch 17 Training Avg. Loss: 5.996 Validation Acc: 0.992
Epoch 18 Training Avg. Loss: 6.028 Validation Acc: 0.992
Epoch 19 Training Avg. Loss: 5.277 Validation Acc: 0.992
Epoch 20 Training Avg. Loss: 4.744 Validation Acc: 0.991
Saving model in ./model/
```

```
[23]: ### Calculate prediction accuracy
      ### on test set
```

```

### restoring the saved model
del g
## create a new graph
## and build the model
g2 = tf.Graph()
with g2.as_default():
    tf.set_random_seed(random_seed)
    ## build the graph
    build_cnn()
    ## saver:
    saver = tf.train.Saver()

## create a new session
## and restore the model
with tf.Session(graph=g2) as sess:
    load(saver, sess, epoch=20, path='./model/')
    preds = predict(sess, X_test_centered, return_proba=False)
    print('Test Accuracy: %.3f%%' % (100*np.sum(preds == y_test)/len(y_test)))

```

Building 1st layer:

```

<tf.Variable 'conv_1/_weights:0' shape=(5, 5, 1, 32) dtype=float32_ref>
<tf.Variable 'conv_1/_biases:0' shape=(32,) dtype=float32_ref>
Tensor("conv_1/Conv2D:0", shape=(?, 24, 24, 32), dtype=float32)
Tensor("conv_1/net_pre-activation:0", shape=(?, 24, 24, 32), dtype=float32)
Tensor("conv_1/activation:0", shape=(?, 24, 24, 32), dtype=float32)

```

Building 2nd layer:

```

<tf.Variable 'conv_2/_weights:0' shape=(5, 5, 32, 64) dtype=float32_ref>
<tf.Variable 'conv_2/_biases:0' shape=(64,) dtype=float32_ref>
Tensor("conv_2/Conv2D:0", shape=(?, 8, 8, 64), dtype=float32)
Tensor("conv_2/net_pre-activation:0", shape=(?, 8, 8, 64), dtype=float32)
Tensor("conv_2/activation:0", shape=(?, 8, 8, 64), dtype=float32)

```

Building 3rd layer:

```

<tf.Variable 'fc_3/_weights:0' shape=(1024, 1024) dtype=float32_ref>
<tf.Variable 'fc_3/_biases:0' shape=(1024,) dtype=float32_ref>
Tensor("fc_3/MatMul:0", shape=(?, 1024), dtype=float32)
Tensor("fc_3/net_pre-activation:0", shape=(?, 1024), dtype=float32)
Tensor("fc_3/activation:0", shape=(?, 1024), dtype=float32)

```

Building 4th layer:

```

<tf.Variable 'fc_4/_weights:0' shape=(1024, 10) dtype=float32_ref>
<tf.Variable 'fc_4/_biases:0' shape=(10,) dtype=float32_ref>
Tensor("fc_4/MatMul:0", shape=(?, 10), dtype=float32)
Tensor("fc_4/net_pre-activation:0", shape=(?, 10), dtype=float32)
Loading model from ./model/
INFO:tensorflow:Restoring parameters from ./model/cnn-model.ckpt-20

```

Test Accuracy: 99.300%

```
[24]: ## run the prediction on  
## some test samples  
np.set_printoptions(precision=2, suppress=True)  
with tf.Session(graph=g2) as sess:  
    load(saver, sess, epoch=20, path='./model/')  
    print(predict(sess, X_test_centered[:10], return_proba=False))  
    print(predict(sess, X_test_centered[:10], return_proba=True))
```

```
Loading model from ./model/  
INFO:tensorflow:Restoring parameters from ./model/cnn-model.ckpt-20  
[7 2 1 0 4 1 4 9 5 9]  
[[0.  0.  0.  0.  0.  0.  0.  1.  0.  0. ]  
 [0.  0.  1.  0.  0.  0.  0.  0.  0.  0. ]  
 [0.  1.  0.  0.  0.  0.  0.  0.  0.  0. ]  
 [1.  0.  0.  0.  0.  0.  0.  0.  0.  0. ]  
 [0.  0.  0.  0.  1.  0.  0.  0.  0.  0. ]  
 [0.  1.  0.  0.  0.  0.  0.  0.  0.  0. ]  
 [0.  0.  0.  0.  1.  0.  0.  0.  0.  0. ]  
 [0.  0.  0.  0.  0.  0.  0.  0.  0.  1. ]  
 [0.  0.  0.  0.  0.  0.98 0.02 0.  0.  0. ]  
 [0.  0.  0.  0.  0.  0.  0.  0.  0.  1. ]]
```

```
[25]: ## continue training for 20 more epochs  
## without re-initializing :: initialize=False  
## create a new session  
## and restore the model  
with tf.Session(graph=g2) as sess:  
    load(saver, sess, epoch=20, path='./model/')  
    train(sess, training_set=(X_train_centered,   
↪ y_train), validation_set=(X_valid_centered,   
↪ y_valid), initialize=False, epochs=20, random_seed=123)  
    save(saver, sess, epoch=40, path='./model/')  
    preds = predict(sess, X_test_centered, return_proba=False)  
    print('Test Accuracy: %.3f%%' % (100*np.sum(preds == y_test)/len(y_test)))
```

```
Loading model from ./model/  
INFO:tensorflow:Restoring parameters from ./model/cnn-model.ckpt-20  
Epoch 01 Training Avg. Loss:  4.276  Validation Acc:  0.990  
Epoch 02 Training Avg. Loss:  4.328  Validation Acc:  0.992  
Epoch 03 Training Avg. Loss:  3.701  Validation Acc:  0.992  
Epoch 04 Training Avg. Loss:  3.705  Validation Acc:  0.992  
Epoch 05 Training Avg. Loss:  3.089  Validation Acc:  0.992  
Epoch 06 Training Avg. Loss:  2.242  Validation Acc:  0.992  
Epoch 07 Training Avg. Loss:  3.105  Validation Acc:  0.992  
Epoch 08 Training Avg. Loss:  2.375  Validation Acc:  0.992  
Epoch 09 Training Avg. Loss:  2.084  Validation Acc:  0.992  
Epoch 10 Training Avg. Loss:  2.455  Validation Acc:  0.993
```

```

Epoch 11 Training Avg. Loss: 2.579 Validation Acc: 0.993
Epoch 12 Training Avg. Loss: 2.050 Validation Acc: 0.992
Epoch 13 Training Avg. Loss: 2.555 Validation Acc: 0.992
Epoch 14 Training Avg. Loss: 1.100 Validation Acc: 0.991
Epoch 15 Training Avg. Loss: 1.336 Validation Acc: 0.992
Epoch 16 Training Avg. Loss: 1.813 Validation Acc: 0.993
Epoch 17 Training Avg. Loss: 1.311 Validation Acc: 0.992
Epoch 18 Training Avg. Loss: 1.477 Validation Acc: 0.991
Epoch 19 Training Avg. Loss: 1.855 Validation Acc: 0.992
Epoch 20 Training Avg. Loss: 1.301 Validation Acc: 0.992
Saving model in ./model/
Test Accuracy: 99.350%

```

```

[26]: import tensorflow as tf
import numpy as np
class ConvNN(object):
    def __init__(self, batchsize=64, epochs=20,
↳ learning_rate=1e-4, dropout_rate=0.5, shuffle=True, random_seed=None):
        np.random.seed(random_seed)
        self.batchsize = batchsize
        self.epochs = epochs
        self.learning_rate = learning_rate
        self.dropout_rate = dropout_rate
        self.shuffle = shuffle
        g = tf.Graph()
        with g.as_default():
            ## set random-seed:
            tf.set_random_seed(random_seed)
            ## build the network:
            self.build()
            ## initializer
            self.init_op = tf.global_variables_initializer()
            ## saver
            self.saver = tf.train.Saver()

            ## create a session
            self.sess = tf.Session(graph=g)

    def build(self):

        ## Placeholders for X and y:
        tf_x = tf.placeholder(tf.float32, shape=[None, 784], name='tf_x')
        tf_y = tf.placeholder(tf.int32, shape=[None], name='tf_y')
        is_train = tf.placeholder(tf.bool, shape=(), name='is_train')
        ## reshape x to a 4D tensor:
        ## [batchsize, width, height, 1]

```

```

        tf_x_image = tf.reshape(tf_x, shape=[-1, 28, 28, 1
↪1], name='input_x_2dimages')
        ## One-hot encoding:
        tf_y_onehot = tf.one_hot(indices=tf_y, depth=10, dtype=tf.
↪float32, name='input_y_onehot')
        ## 1st layer: Conv_1
        h1 = tf.layers.conv2d(tf_x_image, kernel_size=(5, 1
↪5), filters=32, activation=tf.nn.relu)
        ## MaxPooling
        h1_pool = tf.layers.max_pooling2d(h1, pool_size=(2, 2), strides=(2, 2))
        ## 2nd layer: Conv_2
        h2 = tf.layers.conv2d(h1_pool, kernel_size=(5, 1
↪5), filters=64, activation=tf.nn.relu)
        ## MaxPooling
        h2_pool = tf.layers.max_pooling2d(h2, pool_size=(2, 2), strides=(2, 2))
        ## 3rd layer: Fully Connected
        input_shape = h2_pool.get_shape().as_list()
        n_input_units = np.prod(input_shape[1:])
        h2_pool_flat = tf.reshape(h2_pool, shape=[-1, n_input_units])
        h3 = tf.layers.dense(h2_pool_flat, 1024, activation=tf.nn.relu)
        ## Dropout
        h3_drop = tf.layers.dropout(h3, rate=self.dropout_rate, training=is_train)
        ## 4th layer: Fully Connected (linear activation)
        h4 = tf.layers.dense(h3_drop, 10, activation=None)
        ## Prediction
        predictions = {
            'probabilities': tf.nn.softmax(h4, name='probabilities'),
            'labels': tf.cast(tf.argmax(h4, axis=1),
            tf.int32, name='labels')
        }

        ## Loss Function and Optimization
        cross_entropy_loss = tf.reduce_mean(tf.nn.
↪softmax_cross_entropy_with_logits(logits=h4, 1
↪labels=tf_y_onehot), name='cross_entropy_loss')
        ## Optimizer:
        optimizer = tf.train.AdamOptimizer(self.learning_rate)
        optimizer = optimizer.minimize(cross_entropy_loss, name='train_op')
        ## Finding accuracy
        correct_predictions = tf.equal(predictions['labels'], tf_y, 1
↪name='correct_preds')
        accuracy = tf.reduce_mean(tf.cast(correct_predictions, tf.
↪float32), name='accuracy')
        def save(self, epoch, path='./tflayers-model/'):
            if not os.path.isdir(path):
                os.makedirs(path)

```

```

        print('Saving model in %s' % path)
        self.saver.save(self.sess,os.path.join(path, 'model.
↪ckpt'),global_step=epoch)

    def load(self, epoch, path):
        print('Loading model from %s' % path)
        self.saver.restore(self.sess,os.path.join(path, 'model.ckpt-%d' %
↪epoch))

    def train(self, training_set,validation_set=None,initialize=True):
        ## initialize variables
        if initialize:
            self.sess.run(self.init_op)
        self.train_cost_ = []
        X_data = np.array(training_set[0])
        y_data = np.array(training_set[1])
        for epoch in range(1, self.epochs+1):
            batch_gen = batch_generator(X_data, y_data,shuffle=self.shuffle)
            avg_loss = 0.0
            for i, (batch_x,batch_y) in enumerate(batch_gen):
                feed = {'tf_x:0': batch_x,'tf_y:0': batch_y,'is_train:0': True}
↪## for dropout
                loss, _ = self.sess.run(['cross_entropy_loss:0',
↪'train_op'],feed_dict=feed)
                avg_loss += loss

            print('Epoch %02d: Training Avg. Loss: '%7.3f' % (epoch,
↪avg_loss), end=' ')
            if validation_set is not None:
                feed = {'tf_x:0': batch_x,'tf_y:0': batch_y,'is_train:0' :
↪False} ## for dropout
                valid_acc = self.sess.run('accuracy:0',feed_dict=feed)
                print('Validation Acc: %7.3f' % valid_acc)
            else:
                print()

    def predict(self, X_test, return_proba=False):
        feed = {'tf_x:0' : X_test,'is_train:0' : False} ## for dropout
        if return_proba:
            return self.sess.run('probabilities:0',feed_dict=feed)
        else:
            return self.sess.run('labels:0',feed_dict=feed)

```

```

[27]: cnn = ConvNN(random_seed=123)
      ## train the model
      cnn.train(training_set=(X_train_centered,
↪y_train),validation_set=(X_valid_centered, y_valid),initialize=True)

```

```
cnn.save(epoch=20)
```

```
Epoch 01: Training Avg. Loss: 262.236 Validation Acc: 0.938
Epoch 02: Training Avg. Loss: 73.340 Validation Acc: 1.000
Epoch 03: Training Avg. Loss: 50.417 Validation Acc: 1.000
Epoch 04: Training Avg. Loss: 39.019 Validation Acc: 1.000
Epoch 05: Training Avg. Loss: 31.782 Validation Acc: 1.000
Epoch 06: Training Avg. Loss: 26.602 Validation Acc: 1.000
Epoch 07: Training Avg. Loss: 22.542 Validation Acc: 1.000
Epoch 08: Training Avg. Loss: 19.906 Validation Acc: 1.000
Epoch 09: Training Avg. Loss: 17.299 Validation Acc: 1.000
Epoch 10: Training Avg. Loss: 15.416 Validation Acc: 1.000
Epoch 11: Training Avg. Loss: 12.735 Validation Acc: 1.000
Epoch 12: Training Avg. Loss: 11.523 Validation Acc: 1.000
Epoch 13: Training Avg. Loss: 10.074 Validation Acc: 1.000
Epoch 14: Training Avg. Loss: 8.626 Validation Acc: 1.000
Epoch 15: Training Avg. Loss: 8.084 Validation Acc: 1.000
Epoch 16: Training Avg. Loss: 7.189 Validation Acc: 1.000
Epoch 17: Training Avg. Loss: 6.973 Validation Acc: 1.000
Epoch 18: Training Avg. Loss: 5.141 Validation Acc: 1.000
Epoch 19: Training Avg. Loss: 5.054 Validation Acc: 1.000
Epoch 20: Training Avg. Loss: 5.591 Validation Acc: 1.000
Saving model in ./tflayers-model/
```

```
[28]: del cnn
      cnn2 = ConvNN(random_seed=123)
      cnn2.load(epoch=20, path='./tflayers-model/')
      print(cnn2.predict(X_test_centered[:10, :]))
```

```
Loading model from ./tflayers-model/
INFO:tensorflow:Restoring parameters from ./tflayers-model/model.ckpt-20
[7 2 1 0 4 1 4 9 5 9]
```

```
[29]: preds = cnn2.predict(X_test_centered)
      print('Test Accuracy: %.2f%%' % (100*np.sum(y_test == preds)/len(y_test)))
```

```
Test Accuracy: 99.37%
```

```
[ ]:
```